

# Deep Learning in Scientific Computing

## PINNs – applications

Spring Semester 2023

Siddhartha Mishra  
Ben Moseley



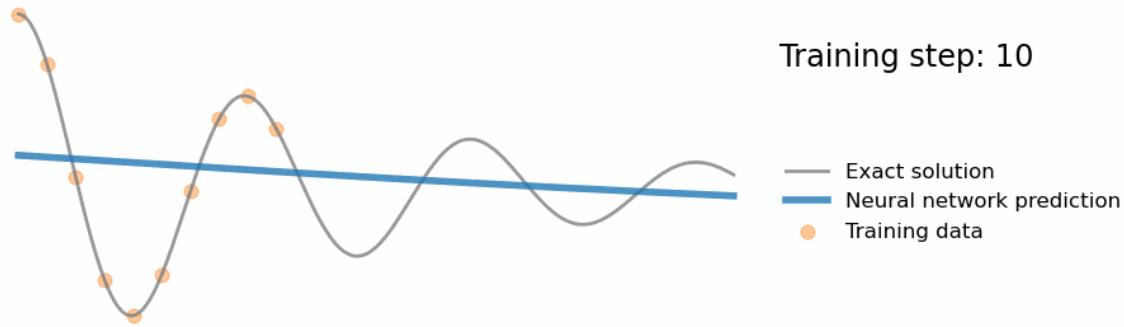
# Lecture overview

- Introduction to physics-informed neural networks (PINNs): recap
- Training PINNs
- Live-coding a PINN in PyTorch
- PINN applications
  - Simulation
  - Inversion
  - Equation discovery

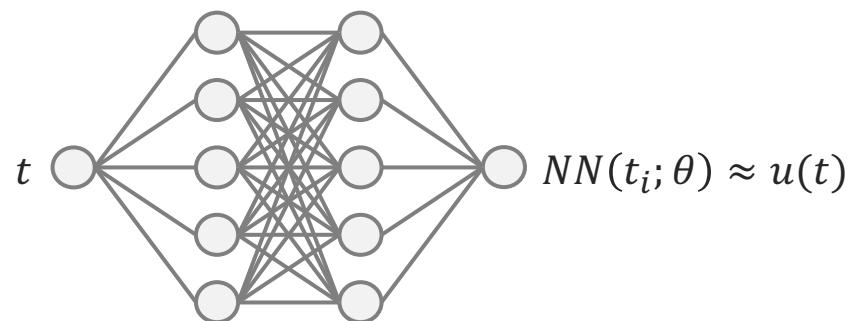
# Course timeline

Tutorials		Lectures
Tue 3:15-14:00, HG E5		Fri 12:15-14:00, HG D1.1
21.02.		24.02. <del>Course introduction</del>
28.02.	<del>Intro to PyTorch</del>	03.03. <del>Introduction to deep learning I</del>
07.03.	<del>Simple DNNs in PyTorch</del>	10.03. <del>Introduction to deep learning II</del>
14.03.	<del>Advanced DNNs in PyTorch</del>	17.03. <del>Physics-informed neural networks – introduction and theory</del>
21.03.	<del>PINN exercises</del>	24.03. <b>Physics-informed neural networks – applications</b>
28.03.	Implementing PINNs I	31.03. Physics-informed neural networks – extensions
04.04.	Implementing PINNs II	07.04.
11.04.		14.04.
18.04.	Introduction to projects	21.04. Neural operators – introduction and theory
25.04.	Implementing neural operators I	28.04. Neural operators – applications
02.05.	Implementing neural operators II	05.05. Neural operators – extensions
09.05.	Operator learning exercises	12.05. Graph and sequence models
16.05.	Project discussions	19.05. Differentiable physics – introduction
23.05.	Implementing autodifferentiation	26.05. Differentiable physics and neural differential equations
30.05.	Intro to JAX	02.06. Future trends and overview of CAMLAB

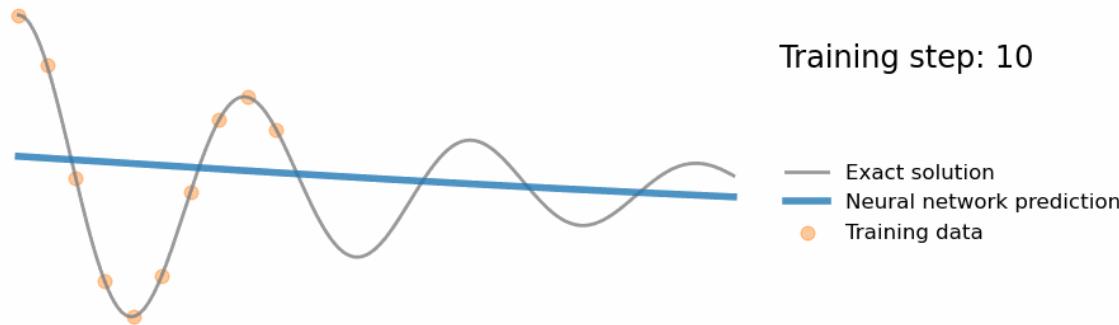
# Naïve neural network



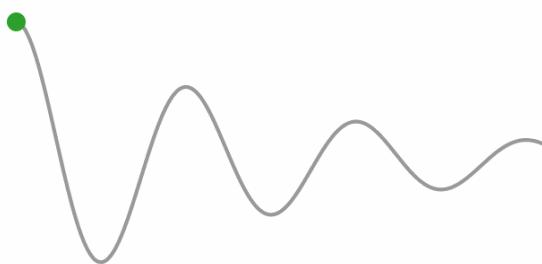
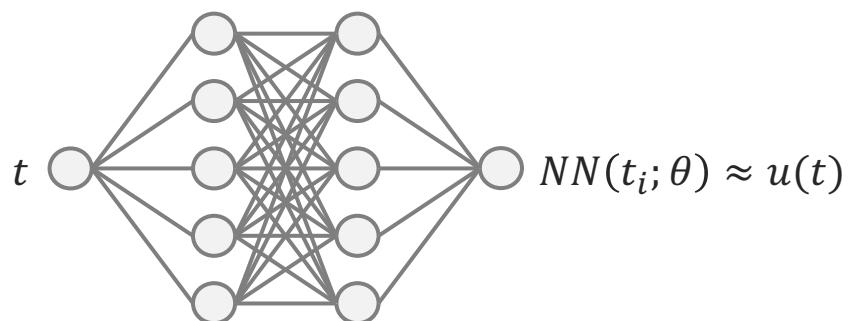
$$L(\theta) = \frac{1}{N} \sum_i^N (NN(t_i; \theta) - \underline{u}_i)^2$$



# Naïve neural network



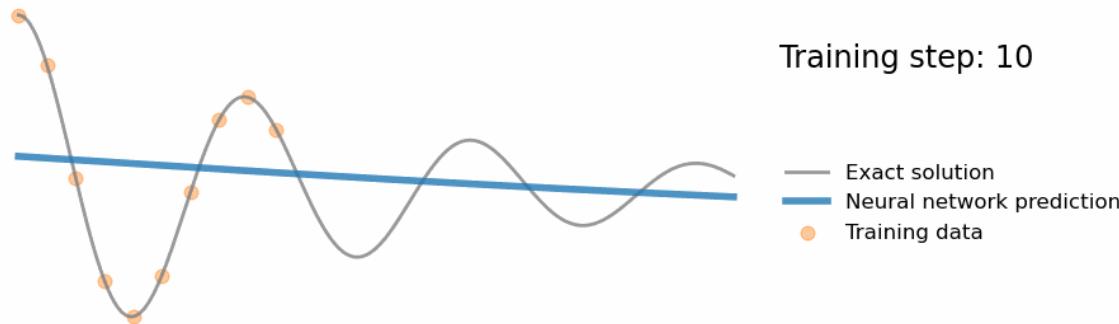
$$L(\theta) = \frac{1}{N} \sum_i^N (NN(t_i; \theta) - u_i)^2$$



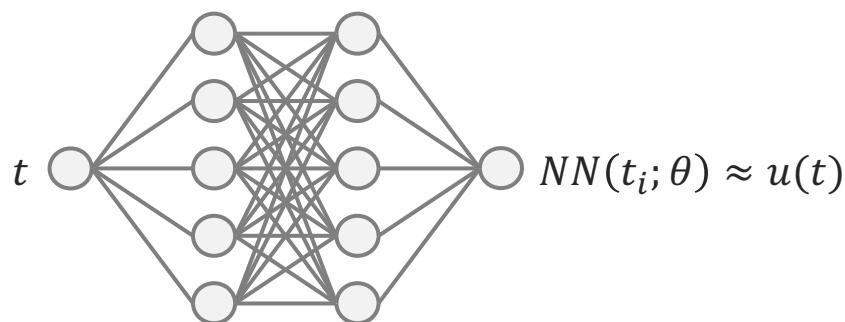
Damped harmonic oscillator

$$m \frac{d^2u}{dt^2} + \mu \frac{du}{dt} + ku = 0$$

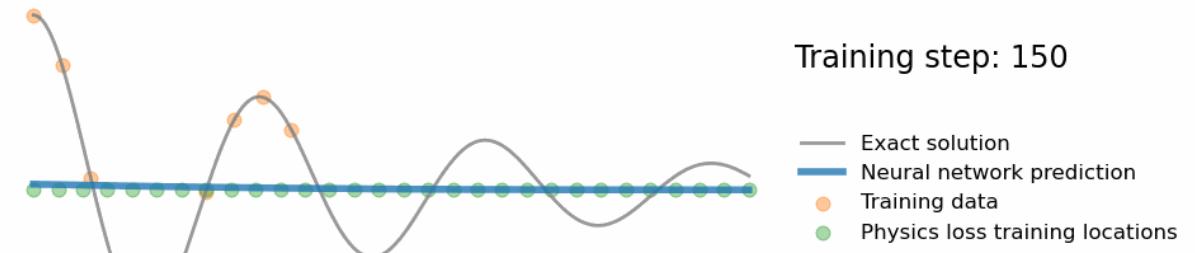
## Naïve neural network



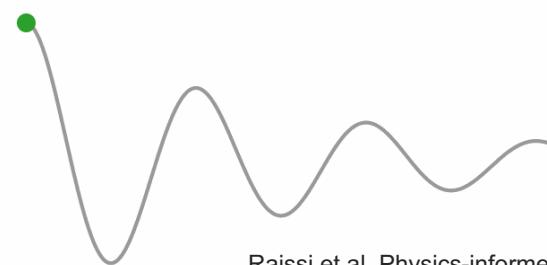
$$L(\theta) = \frac{1}{N} \sum_i^N (NN(t_i; \theta) - u_i)^2$$



## Physics-informed neural network



$$L(\theta) = \frac{1}{N} \sum_i^N (NN(t_i; \theta) - u_i)^2 + \frac{\lambda}{M} \sum_j^M \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_j; \theta) \right)^2$$

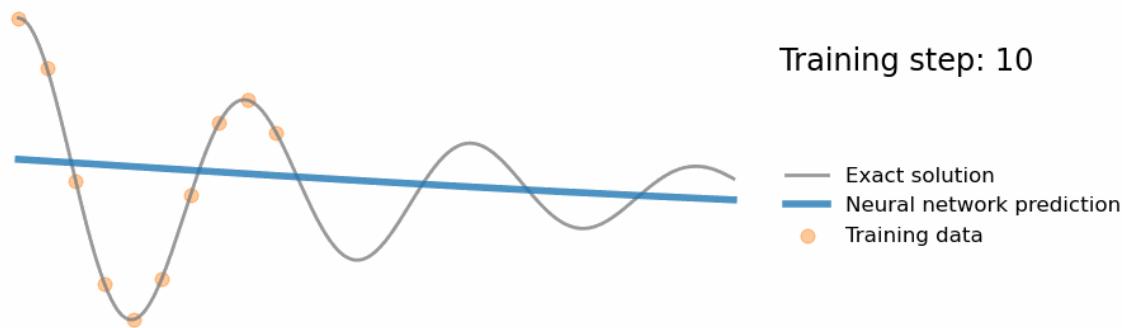


Damped harmonic oscillator

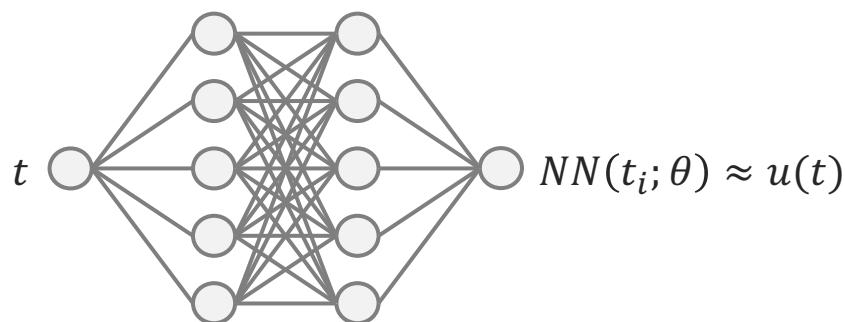
$$m \frac{d^2u}{dt^2} + \mu \frac{du}{dt} + ku = 0$$

Raissi et al, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, JCP (2018)  
 Lagaris et al, Artificial neural networks for solving ordinary and partial differential equations, IEEE (1998)

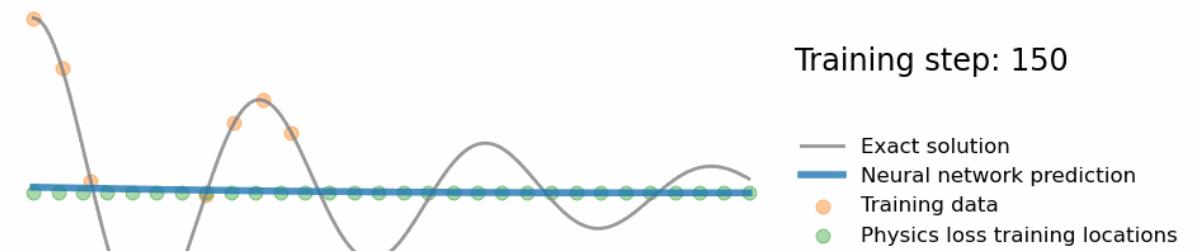
## Naïve neural network



$$L(\theta) = \frac{1}{N} \sum_i^N (NN(t_i; \theta) - u_i)^2$$

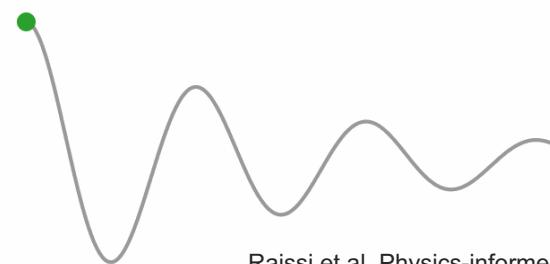


## Physics-informed neural network



$$L(\theta) = \frac{1}{N} \sum_i^N (NN(t_i; \theta) - u_i)^2 \quad \text{Supervised loss}$$

$$+ \frac{\lambda}{M} \sum_j^M \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_j; \theta) \right)^2 \quad \text{Physics loss aka PDE residual}$$

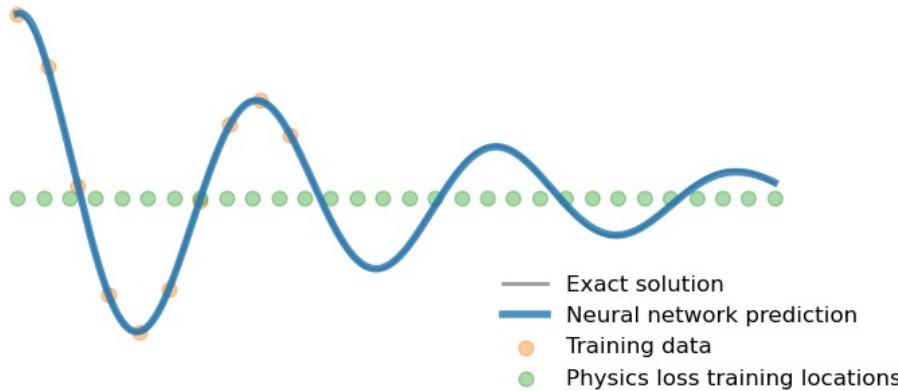


Damped harmonic oscillator

$$m \frac{d^2u}{dt^2} + \mu \frac{du}{dt} + ku = 0$$

Raissi et al, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, JCP (2018)  
 Lagaris et al, Artificial neural networks for solving ordinary and partial differential equations, IEEE (1998)

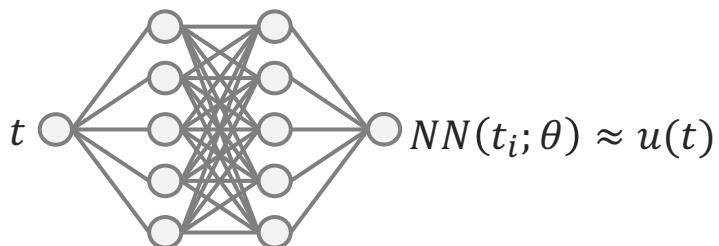
# Physics-informed neural network (PINN)



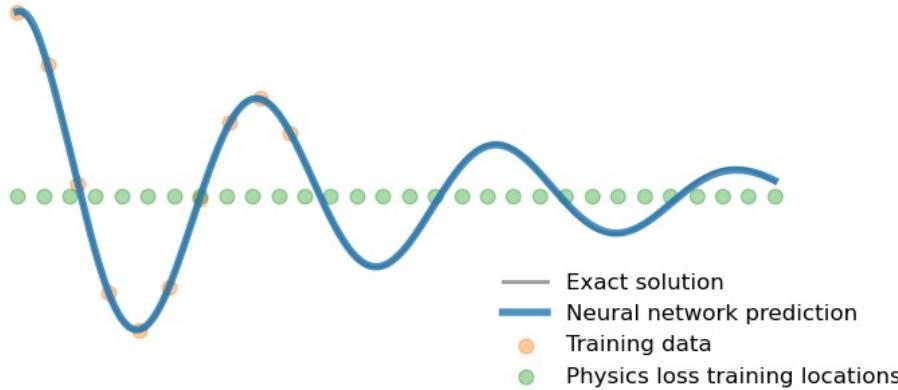
From a ML perspective:

- Physics loss is an **unsupervised** regulariser, which adds prior knowledge

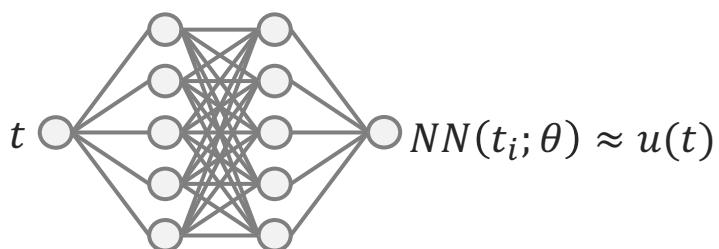
$$L(\theta) = \frac{1}{N} \sum_i^N (NN(t_i; \theta) - u_i)^2 \quad \text{Supervised loss}$$
$$+ \frac{\lambda}{M} \sum_j^M \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_j; \theta) \right)^2 \quad \text{Physics loss aka PDE residual}$$



# Physics-informed neural network (PINN)



$$L(\theta) = \frac{1}{N} \sum_i^N (NN(t_i; \theta) - u_i)^2 \quad \text{Supervised loss}$$
$$+ \frac{\lambda}{M} \sum_j^M \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_j; \theta) \right)^2 \quad \text{Physics loss aka PDE residual}$$



From a ML perspective:

- Physics loss is an **unsupervised** regulariser, which adds prior knowledge

From a mathematical perspective:

- PINNs provide a way to solve PDEs:
  - Neural network is a mesh-free, **functional** approximation of PDE solution
  - Physics loss is used to assert solution is **consistent** with PDE
  - Supervised loss is used to assert boundary/initial conditions, to ensure solution is **unique**

# PINNs are a general framework for solving PDEs

Given a PDE and its boundary/initial conditions

$$\begin{aligned}\mathcal{D}[u(x)] &= f(x), \quad x \in \Omega \subset \mathbb{R}^d \\ \mathcal{B}_k[u(x)] &= g_k(x), \quad x \in \Gamma_k \subset \partial\Omega\end{aligned}$$

Where  $\mathcal{D}$  is some differential operator,  $\mathcal{B}_k$  are a set of boundary operators, and  $u(x)$  is the solution to the PDE

PINNs train a neural network to **approximate** the solution to the PDE  $NN(x; \theta) \approx u(x)$  using the following loss function:

$$L(\theta) = L_b(\theta) + L_p(\theta)$$

$$L_b(\theta) = \sum_k \frac{\lambda_k}{N_{bk}} \sum_j^{|N_{bk}|} \| \mathcal{B}_k[NN(x_{kj}; \theta)] - g_k(x_{kj}) \|^2 \text{ **Boundary loss**}$$

$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \| \mathcal{D}[NN(x_i; \theta)] - f(x_i) \|^2 \text{ **Physics loss**}$$

# PINNs are a general framework for solving PDEs

Given a PDE and its boundary/initial conditions

$$\begin{aligned}\mathcal{D}[u(x)] &= f(x), \quad x \in \Omega \subset \mathbb{R}^d \\ \mathcal{B}_k[u(x)] &= g_k(x), \quad x \in \Gamma_k \subset \partial\Omega\end{aligned}$$

Where  $\mathcal{D}$  is some differential operator,  $\mathcal{B}_k$  are a set of boundary operators, and  $u(x)$  is the solution to the PDE

PINNs train a neural network to **approximate** the solution to the PDE  $NN(x; \theta) \approx u(x)$  using the following loss function:

$$L(\theta) = L_b(\theta) + L_p(\theta)$$

$$L_b(\theta) = \sum_k \frac{\lambda_k}{N_{bk}} \sum_j^{N_{bk}} \| \mathcal{B}_k[NN(x_{kj}; \theta)] - g_k(x_{kj}) \|^2 \quad \textbf{Boundary loss}$$

$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \| \mathcal{D}[NN(x_i; \theta)] - f(x_i) \|^2 \quad \textbf{Physics loss}$$

For example, the 1+1D viscous Burgers' equation:

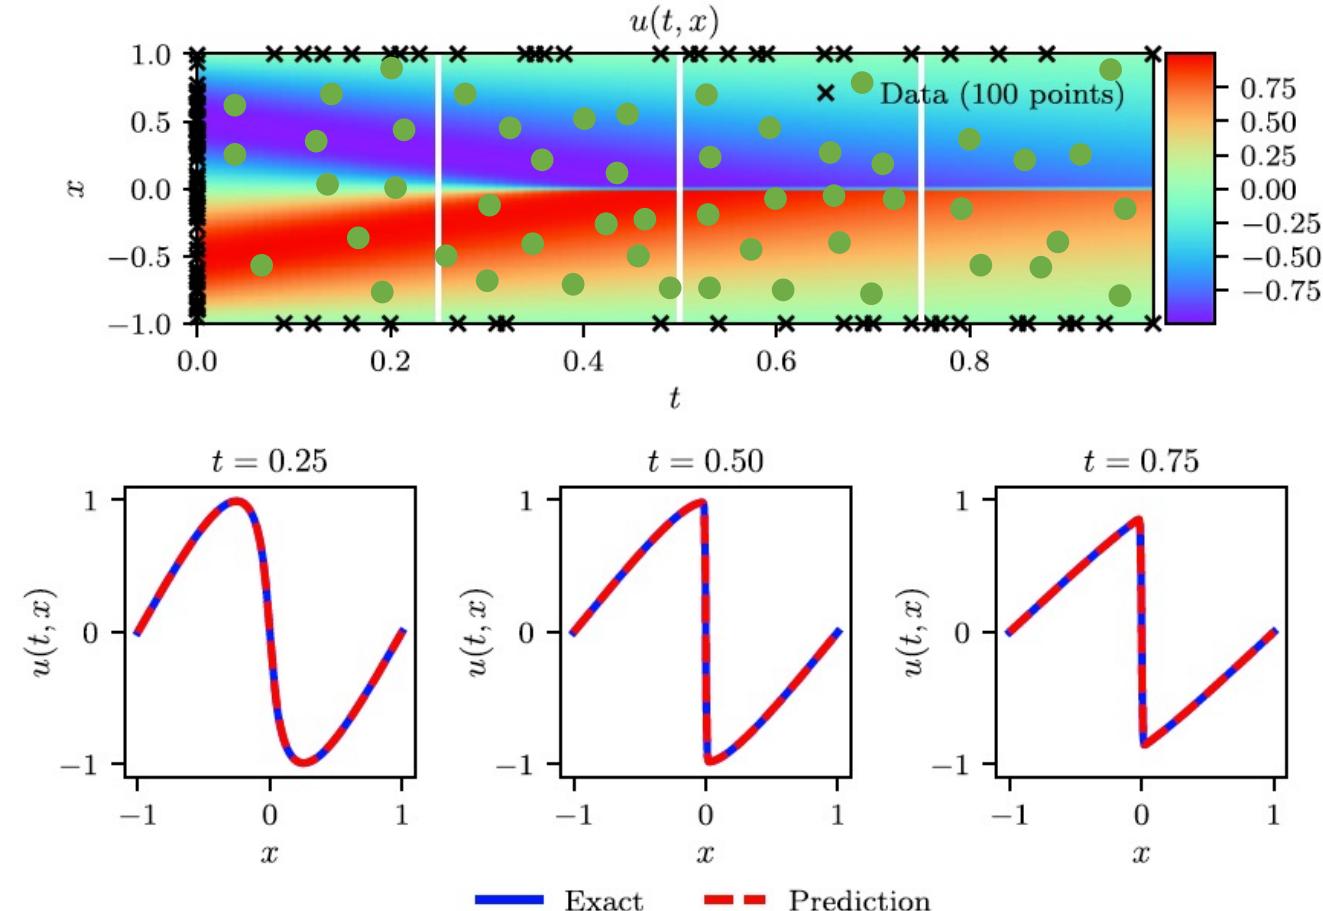
$$\frac{\partial u}{\partial t} + u \frac{\partial u}{\partial x} - \nu \frac{\partial^2 u}{\partial x^2} = 0$$

$$\begin{aligned}u(x, 0) &= -\sin(\pi x) \\ u(-1, t) &= u(+1, t) = 0\end{aligned}$$

$$NN(x, t; \theta) \approx u(x, t)$$

$$\begin{aligned}L_b(\theta) &= \frac{\lambda_1}{N_{b1}} \sum_j^{N_{b1}} (NN(x_j, 0; \theta) + \sin(\pi x_j))^2 \\ &\quad + \frac{\lambda_2}{N_{b2}} \sum_k^{N_{b2}} (NN(-1, t_k; \theta) - 0)^2 \\ &\quad + \frac{\lambda_3}{N_{b3}} \sum_l^{N_{b3}} (NN(+1, t_l; \theta) - 0)^2 \\ L_p(\theta) &= \frac{1}{N_p} \sum_i^{N_p} \left( \left( \frac{\partial NN}{\partial t} + NN \frac{\partial NN}{\partial x} - \nu \frac{\partial^2 NN}{\partial x^2} \right)(x_i, t_i; \theta) \right)^2\end{aligned}$$

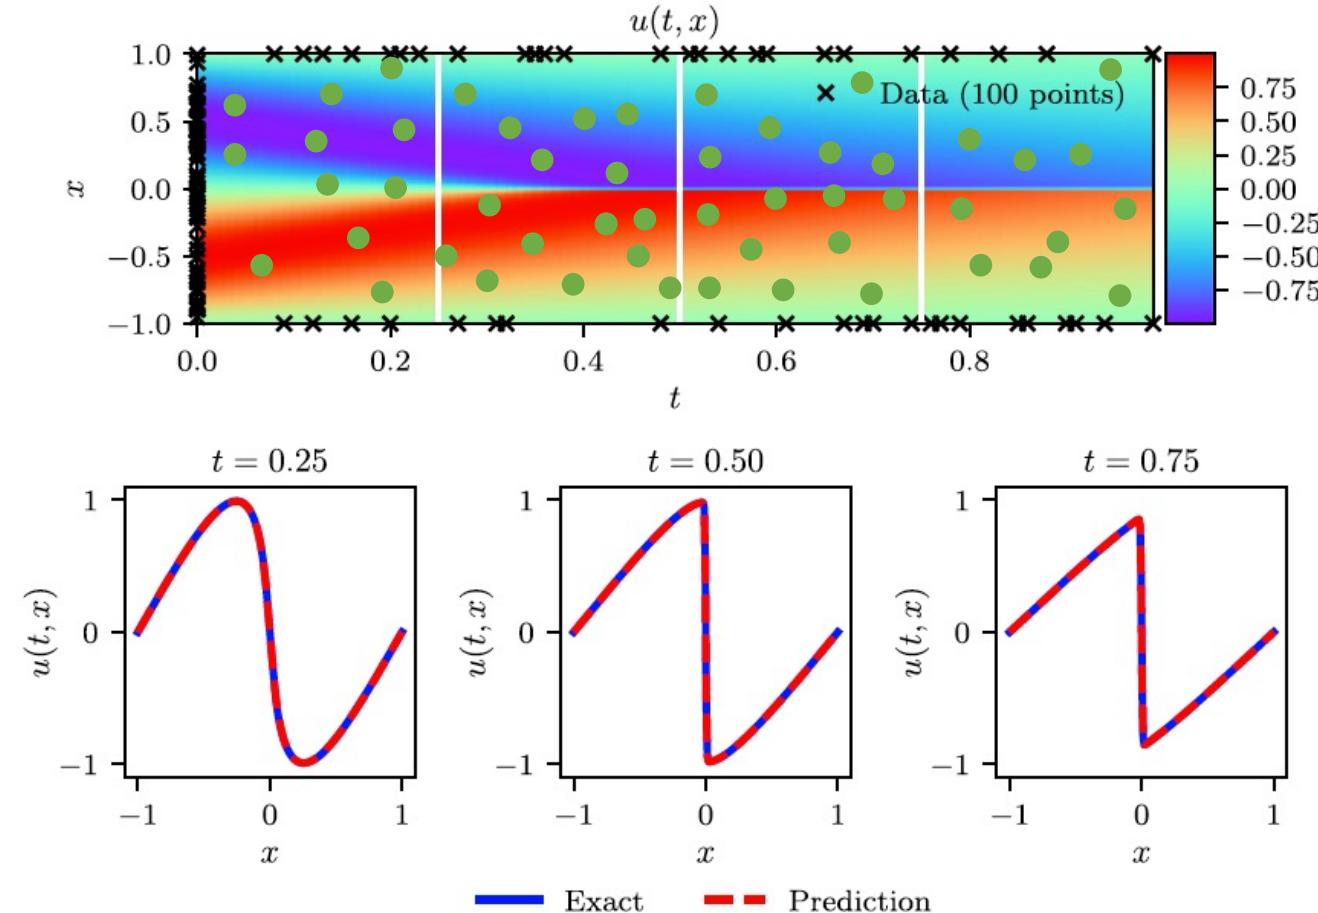
# PINNs for solving viscous Burgers' equation



$$L_b(\theta) = \frac{\lambda_1}{N_{b1}} \sum_j^{N_{b1}} (NN(x_j, 0; \theta) + \sin(\pi x_j))^2$$
$$+ \frac{\lambda_2}{N_{b2}} \sum_k^{N_{b2}} (NN(-1, t_k; \theta) - 0)^2$$
$$+ \frac{\lambda_3}{N_{b3}} \sum_l^{N_{b3}} (NN(+1, t_l; \theta) - 0)^2$$
$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \left( \left( \frac{\partial NN}{\partial t} + NN \frac{\partial NN}{\partial x} - \nu \frac{\partial^2 NN}{\partial x^2} \right) (x_i, t_i; \theta) \right)^2$$

Raissi et al, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, JCP (2018)

# PINNs for solving viscous Burgers' equation



Raissi et al, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, JCP (2018)

$$L_b(\theta) = \frac{\lambda_1}{N_{b1}} \sum_j^{N_{b1}} (NN(x_j, 0; \theta) + \sin(\pi x_j))^2$$
$$+ \frac{\lambda_2}{N_{b2}} \sum_k^{N_{b2}} (NN(-1, t_k; \theta) - 0)^2$$
$$+ \frac{\lambda_3}{N_{b3}} \sum_l^{N_{b3}} (NN(+1, t_l; \theta) - 0)^2$$
$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \left( \left( \frac{\partial NN}{\partial t} + NN \frac{\partial NN}{\partial x} - \nu \frac{\partial^2 NN}{\partial x^2} \right) (x_i, t_i; \theta) \right)^2$$

$$\nu = 0.01/\pi$$

$N_p = 10,000$  (Latin hypercube sampling)

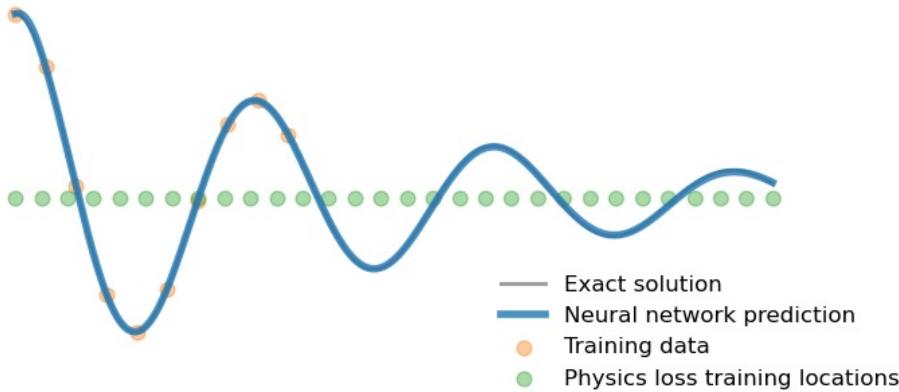
$$N_{b1} + N_{b2} + N_{b3} = 100$$

Fully connected network with 9 layers, 20 hidden units (3021 free parameters)

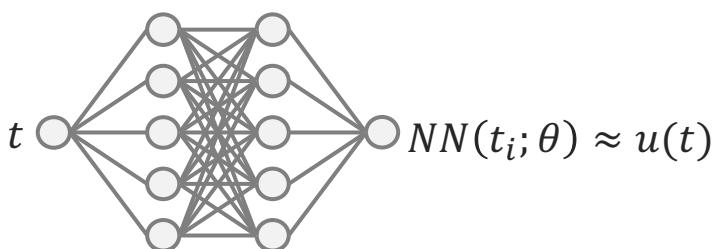
Tanh activation function

L-BFGS optimiser

# PINN training loop



$$L(\theta) = \frac{1}{N} \sum_i^N (NN(t_i; \theta) - u_i)^2 \quad \text{Supervised loss}$$
$$+ \frac{\lambda}{M} \sum_j^M \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_j; \theta) \right)^2 \quad \text{Physics loss}$$



Training loop:

1. Sample boundary/ physics training points
2. Compute network outputs
3. Compute 1<sup>st</sup> and 2<sup>nd</sup> order gradient of network output **with respect to network input**
4. Compute loss
5. Compute gradient of loss function **with respect to network parameters**
6. Take gradient descent step

How can we compute the gradients (e.g.  $\frac{dNN}{dt}$  and  $\frac{dL}{d\theta}$ ) required in steps 3 and 5?

# Gradient computation for PINNs

Assume network is a MLP, e.g.:

$$NN(t; \theta) = W_3(\sigma(W_2\sigma(W_1t + \mathbf{b}_1) + \mathbf{b}_2) + b_3) \approx u(x)$$

Analogous to backpropagation, we can use the multivariate **chain rule** to compute gradients with respect to network **inputs**:

$$NN(t; \theta) = f \circ \mathbf{g} \circ \mathbf{h}(t; \theta)$$

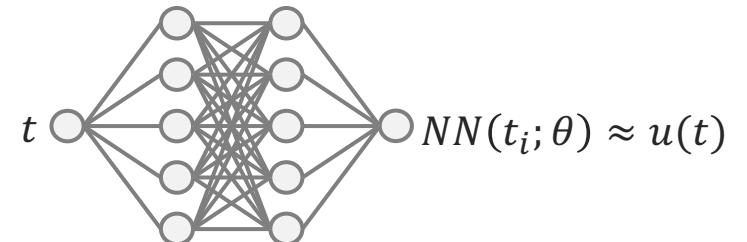
$$f = W_3\sigma(\mathbf{g}) + b_3$$

$$\mathbf{g} = W_2\sigma(\mathbf{h}) + \mathbf{b}_2$$

$$\mathbf{h} = W_1t + \mathbf{b}_1$$

$$\frac{\partial NN}{\partial t} = \frac{\partial f}{\partial \mathbf{g}} \frac{\partial \mathbf{g}}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial t}$$

$$\frac{\partial NN}{\partial t} = W_3 \text{ diag}(\sigma'(\mathbf{g})) W_2 \text{ diag}(\sigma'(\mathbf{h})) W_1$$



(see Lecture 2: Introduction to Deep Learning Part I)

# Extending computational graph



Computing gradients can simply be thought of as **extending** the network's computational graph:

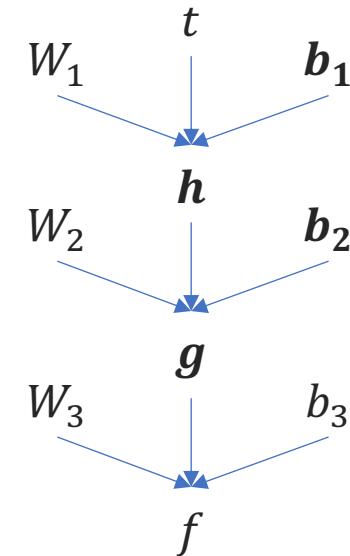
$$NN(t; \theta) = W_3(\sigma(W_2\sigma(W_1t + \mathbf{b}_1) + \mathbf{b}_2) + b_3$$

$$f = W_3\sigma(\mathbf{g}) + b_3$$

$$\mathbf{g} = W_2\sigma(\mathbf{h}) + \mathbf{b}_2$$

$$\mathbf{h} = W_1t + \mathbf{b}_1$$

$$\frac{\partial NN}{\partial t} = W_3 \text{ diag}(\sigma'(\mathbf{g}))W_2 \text{ diag}(\sigma'(\mathbf{h}))W_1$$



# Extending computational graph



Computing gradients can simply be thought of as **extending** the network's computational graph:

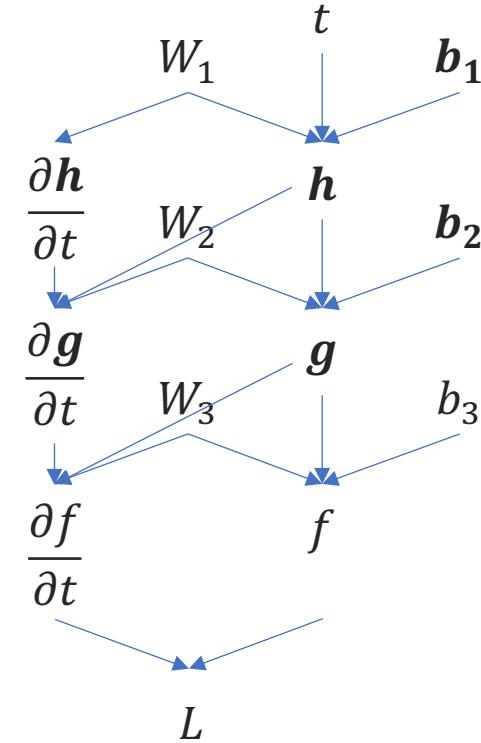
$$NN(t; \theta) = W_3(\sigma(W_2\sigma(W_1t + \mathbf{b}_1) + \mathbf{b}_2) + b_3)$$

$$f = W_3\sigma(\mathbf{g}) + b_3$$

$$\mathbf{g} = W_2\sigma(\mathbf{h}) + \mathbf{b}_2$$

$$\mathbf{h} = W_1t + \mathbf{b}_1$$

$$\frac{\partial NN}{\partial t} = W_3 \text{ diag}(\sigma'(\mathbf{g}))W_2 \text{ diag}(\sigma'(\mathbf{h}))W_1$$



(Forward-mode differentiation)

# Extending computational graph

💡 Computing gradients can simply be thought of as **extending** the network's computational graph:

$$NN(t; \theta) = W_3(\sigma(W_2\sigma(W_1t + \mathbf{b}_1) + \mathbf{b}_2) + b_3)$$

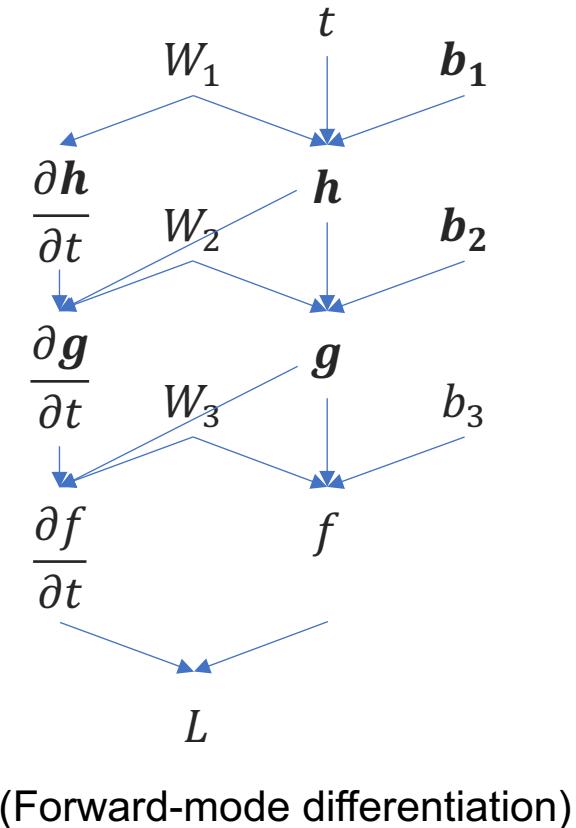
$$f = W_3\sigma(\mathbf{g}) + b_3$$

$$\mathbf{g} = W_2\sigma(\mathbf{h}) + \mathbf{b}_2$$

$$\mathbf{h} = W_1t + \mathbf{b}_1$$

$$\frac{\partial NN}{\partial t} = W_3 \text{ diag}(\sigma'(\mathbf{g}))W_2 \text{ diag}(\sigma'(\mathbf{h}))W_1$$

💡 We can recursively apply **autodifferentiation** to compute gradients and extend the graph



# PINN training loop

Training loop:

1. Sample boundary/ physics training points
2. Compute network outputs
3. Compute 1<sup>st</sup> and 2<sup>nd</sup> order gradient of network output **with respect to network input <= (recursively) apply autodiff, extending graph**
4. Compute loss
5. Compute gradient of loss function **with respect to network parameters <= apply autodiff on extended graph**
6. Take gradient descent step

```
# PINN training psuedocode

#2.
t.requires_grad_(True)# tells PyTorch to start tracking graph
theta.requires_grad_(True)
u = NN(t, theta)

#3.
dudt = torch.autograd.grad(u, t, torch.ones_like(u),
                           create_graph=True)[0]
d2udt2 = torch.autograd.grad(dudt, t, torch.ones_like(u),
                           create_graph=True)[0]

#4.
physics_loss = torch.mean((m*d2udt2 + mu*dudt + k*u)**2)
loss = physics_loss + lambda_*boundary_loss

#5.
dtheta = torch.autograd.grad(loss, theta)[0]
```

 We can recursively apply **autodifferentiation** to compute gradients and extend the graph

# Live-coding a PINN in PyTorch

# Computational cost of higher order derivatives

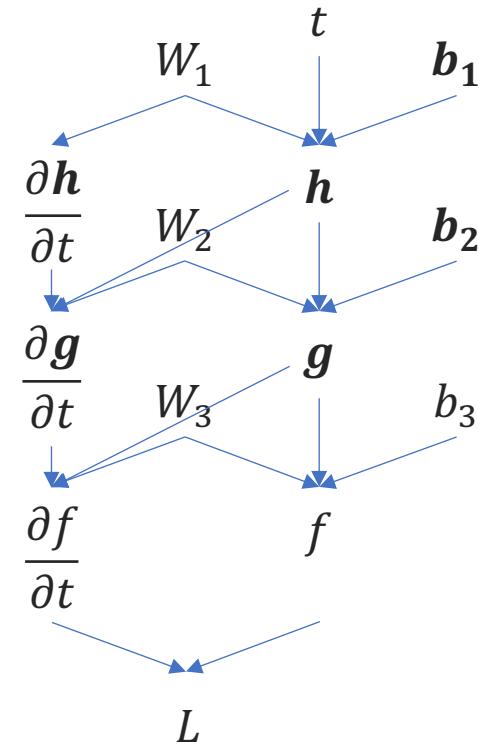
- ⚠ Note, gradient computation roughly doubles the size of the computational graph\*
- ⇒ The cost of evaluating  $\frac{\partial^n f}{\partial t^n}$  grows exponentially with  $n$  (!)
- ⇒ Most time is spent computing gradients, not the forward pass, when training PINNs

\*More precisely, given some  $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$ , its Jacobian  $J \in \mathbb{R}^{m \times n}$  and some vector  $v \in \mathbb{R}^n$

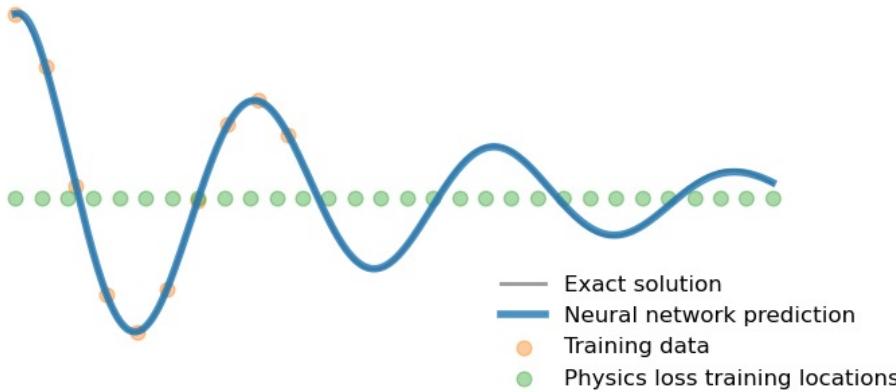
$$TIME\{f, Jv\} \leq \omega TIME\{f\}$$

With a constant  $\omega \in [2, 2.5]$  using autodifferentiation

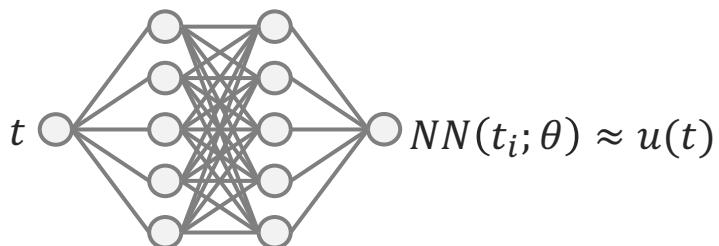
For detailed derivation, see eg: Griewank and Walther, Evaluating Derivatives, Ch 3.1, SIAM (2008))



# Other important training considerations



$$L(\theta) = \frac{1}{N} \sum_i^N (NN(t_i; \theta) - u_i)^2 \quad \text{Supervised loss}$$
$$+ \frac{\lambda}{M} \sum_j^M \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_j; \theta) \right)^2 \quad \text{Physics loss}$$



- $\lambda$  is added as a hyperparameter to control the balance of each loss term (note: units of each loss term are not the same!)
- “Enough” collocation points  $t_j$  must be chosen so that the learned solution is accurate across the full domain
- Similarly, “enough” boundary points must be chosen such that the learned solution is unique
- Training points are usually uniformly, randomly or quasirandomly sampled
- PINNs still suffer from approximation, estimation and optimisation error (just like normal neural networks, see Lecture 2: Introduction to Deep Learning Part II)!

# Lecture overview

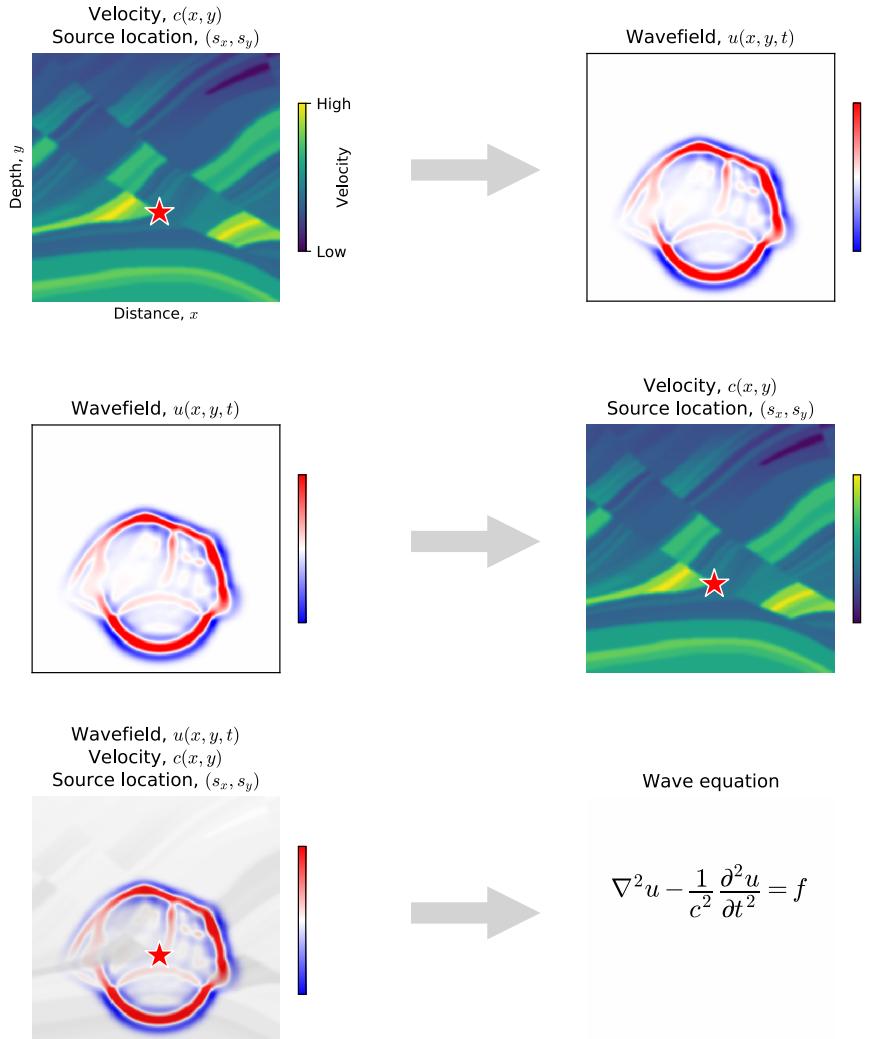
- Introduction to physics-informed neural networks (PINNs): recap
- Training PINNs
- Live-coding a PINN in PyTorch
- PINN applications
  - Simulation
  - Inversion
  - Equation discovery

# 5 min break

# Lecture overview

- Introduction to physics-informed neural networks (PINNs): recap
- Training PINNs
- Live-coding a PINN in PyTorch
- PINN applications
  - Simulation
  - Inversion
  - Equation discovery

# PINN applications



Forward simulation

Estimate wavefield  $u(x, y, t)$   
Given velocity  $c(x, y)$   
and source location  $(s_x, s_y)$

$$b = F(a)$$

Inversion

Estimate velocity  $c(x, y)$   
and source location  $(s_x, s_y)$   
Given wavefield  $u(x, y, t)$

$$b = F(\textcolor{brown}{a})$$

Equation discovery

Estimate governing equation  
Given wavefield  $u(x, y, t)$ ,  
velocity  $c(x, y)$ ,  
and source location  $(s_x, s_y)$

$$b = \textcolor{brown}{F}(a)$$

PINNs can be used to solve **forward**, **inverse** and **equation discovery** tasks related to PDEs

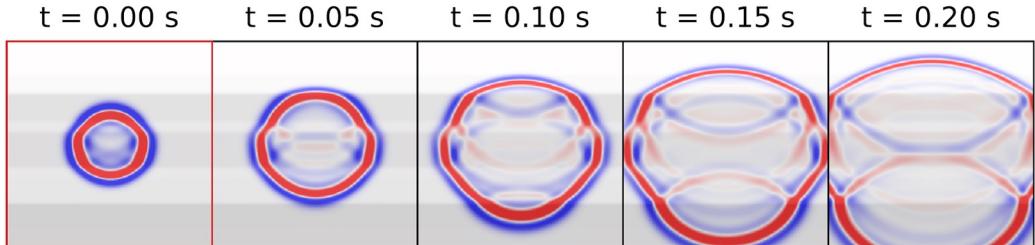
$a$  = set of input conditions

$F$  = physical model of the system

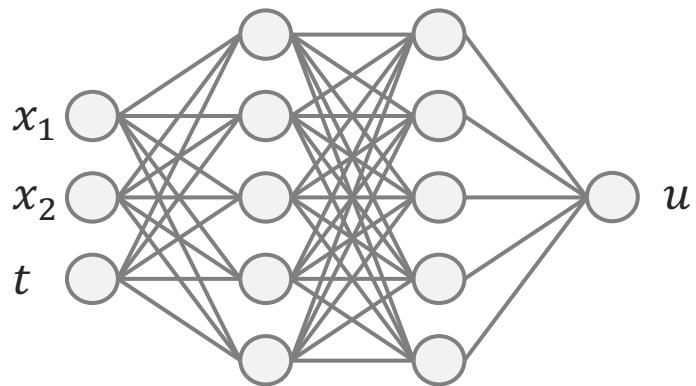
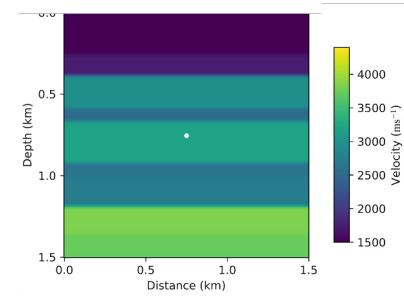
$b$  = resulting properties given  $F$  and  $a$

# PINNs for solving wave equation

Ground truth FD simulation



Velocity model,  $c(x)$



Moseley et al, Solving the wave equation with physics-informed deep learning, ArXiv (2020)

$$L_b(\theta) = \frac{\lambda}{N_b} \sum_j^{N_b} \left( \text{NN}(x_j, t_j; \theta) - \underline{u_{FD}(x_j, t_j)} \right)^2$$

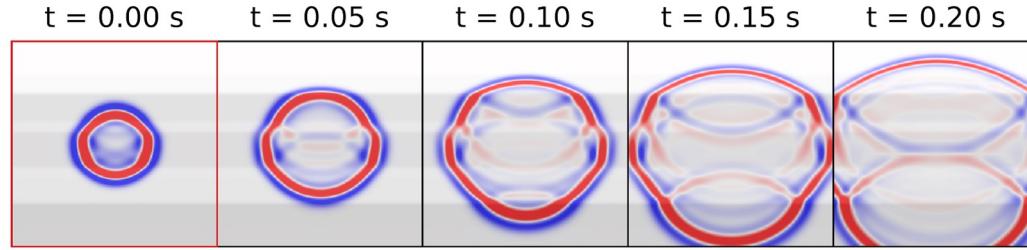
Boundary data from FD simulation (first 0.02 seconds)

$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \left( \left[ \nabla^2 - \frac{1}{c(x_i)^2} \frac{\partial^2}{\partial t^2} \right] \text{NN}(\underline{x_i}, \underline{t_i}; \theta) \right)^2$$

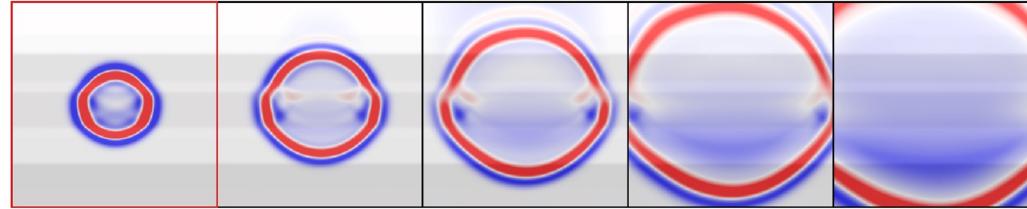
Collocation points randomly sampled over entire domain (up to 0.2 seconds)

# PINNs for solving wave equation

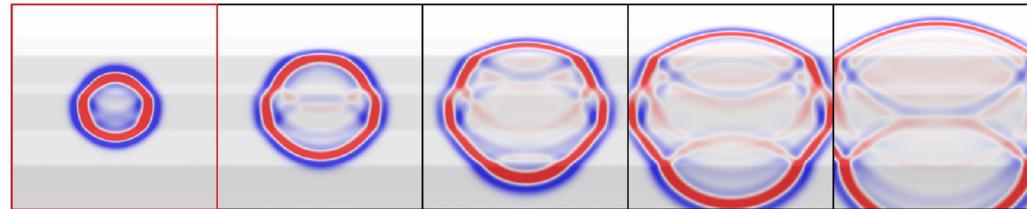
Ground truth FD simulation



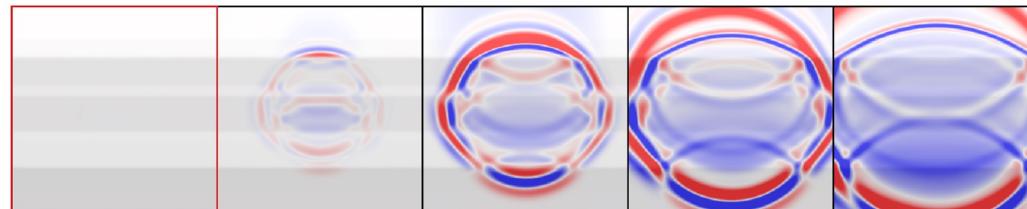
“Naïve” NN



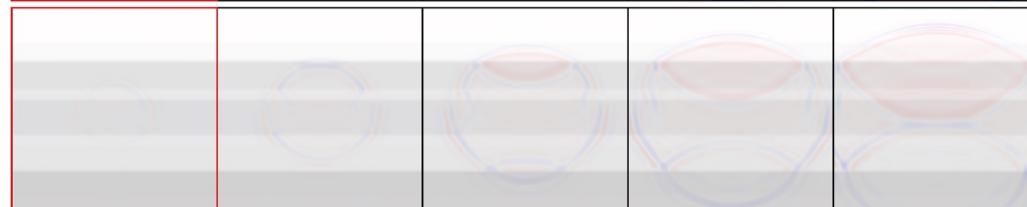
PINN



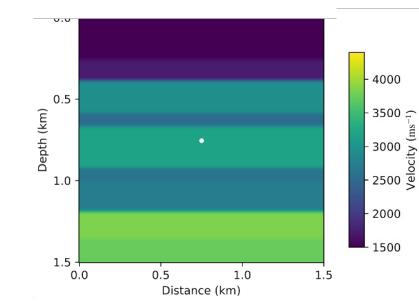
Difference (NN)



Difference (PINN)



Velocity model,  $c(x)$



Moseley et al, Solving the wave equation with physics-informed deep learning, ArXiv (2020)

Mini-batch size  $N_b = N_p = 500$  (random sampling)

Fully connected network with 10 layers, 1024 hidden units

Softplus activation

Adam optimiser

Training time: ~1 hour

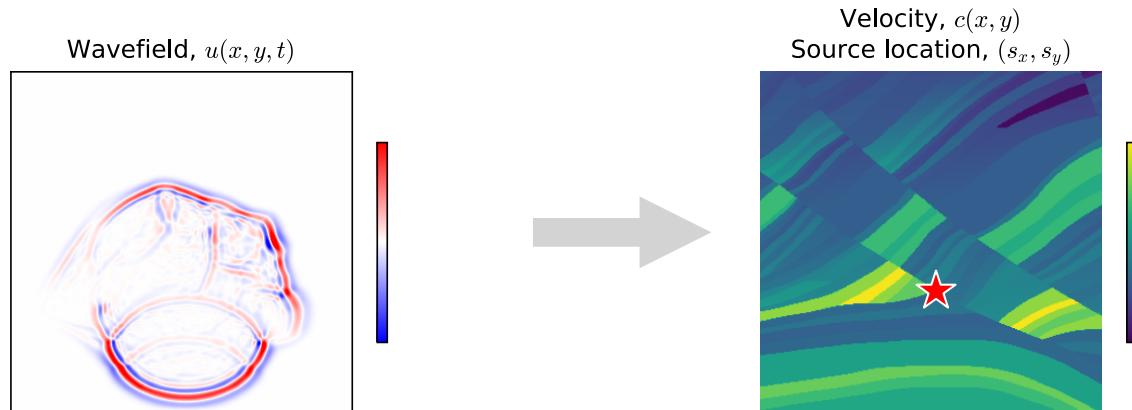
# What is an inverse problem?

Wave equation:

$$\nabla^2 u - \frac{1}{c(x)^2} \frac{\partial^2 u}{\partial t^2} = 0$$

- Fundamentally, inverse problems are **search** problem
- It is often useful to frame them as an optimisation problem, for example:

$$\min_{\hat{a}} \|b - F(\hat{a})\|^2$$



$b$  = observed  
wavefield  $u(x, t)$

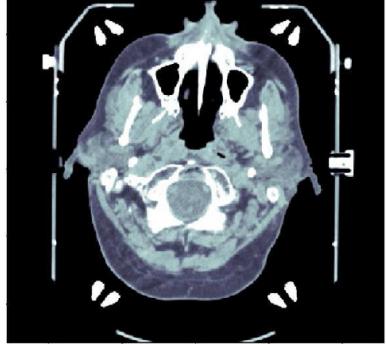
$$b = F(a)$$

$a$  = set of input conditions

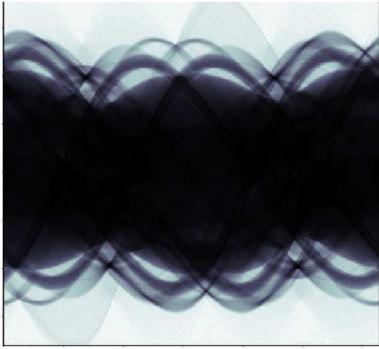
$F$  = physical model of the system

$b$  = resulting properties given  $F$  and  $a$

# Challenges of inverse problems



Ground truth computed tomography image



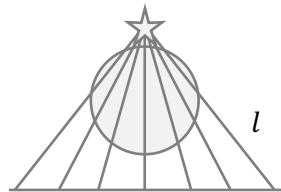
Resulting tomographic data (sinogram)



Result of inverse algorithm (filtered back-projection)

$$a(x)$$

$$F(a)(l) = I_0 \exp\left(-\int_l a(x) dx\right)$$



Adler et al, Solving ill-posed inverse problems using iterative deep neural networks, Inverse Problems (2017)

Typically, inverse problems are **incredibly challenging** to solve because:

- They are usually **ill-posed** (not enough information for a unique solution)
- Observed real-world data is usually **noisy** and **sparse**
- Often require forward modelling to be carried out thousands of times – making them extremely **computationally demanding**

# PINNs for inversion

PINNs for solving **forward** simulation:

$$L(\theta) = L_b(\theta) + L_p(\theta)$$

$$L_b(\theta) = \sum_k \frac{\lambda_k}{N_{bk}} \sum_j^{N_{bk}} \| \mathcal{B}_k [NN(x_{kj}; \theta)] - g_k(x_{kj}) \|^2 \text{ Boundary loss}$$

$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \| \mathcal{D}[NN(x_i; \theta)] - f(x_i) \|^2 \text{ Physics loss}$$

For example:

$$\begin{aligned} D &= \left[ \nabla^2 - \frac{1}{c(x)^2} \frac{\partial^2}{\partial t^2} \right] \\ f &= 0 \end{aligned}$$

# PINNs for inversion

PINNs for solving **forward** simulation:

$$L(\theta) = L_b(\theta) + L_p(\theta)$$

$$L_b(\theta) = \sum_k \frac{\lambda_k}{N_{bk}} \sum_j^{N_{bk}} \|B_k[NN(x_{kj}; \theta)] - g_k(x_{kj})\|^2 \text{ Boundary loss}$$

$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \|\mathcal{D}[NN(x_i; \theta)] - f(x_i)\|^2 \text{ Physics loss}$$

For example:

$$\begin{aligned} D &= \left[ \nabla^2 - \frac{1}{c(x)^2} \frac{\partial^2}{\partial t^2} \right] \\ f &= 0 \end{aligned}$$

PINNs for solving **inverse** problems:

$$L(\theta, \phi) = L_p(\theta, \phi) + L_d(\theta)$$

$$\begin{aligned} L_p(\theta, \phi) &= \frac{1}{N_p} \sum_i^{N_p} \|\mathcal{D}[NN(x_i; \theta); \phi] - f(x_i)\|^2 \text{ Physics loss} \\ L_d(\theta) &= \frac{\lambda}{N_d} \sum_l^{N_d} \|NN(\underline{x}_l; \theta) - \underline{u}_l\|^2 \text{ Data loss} \end{aligned}$$

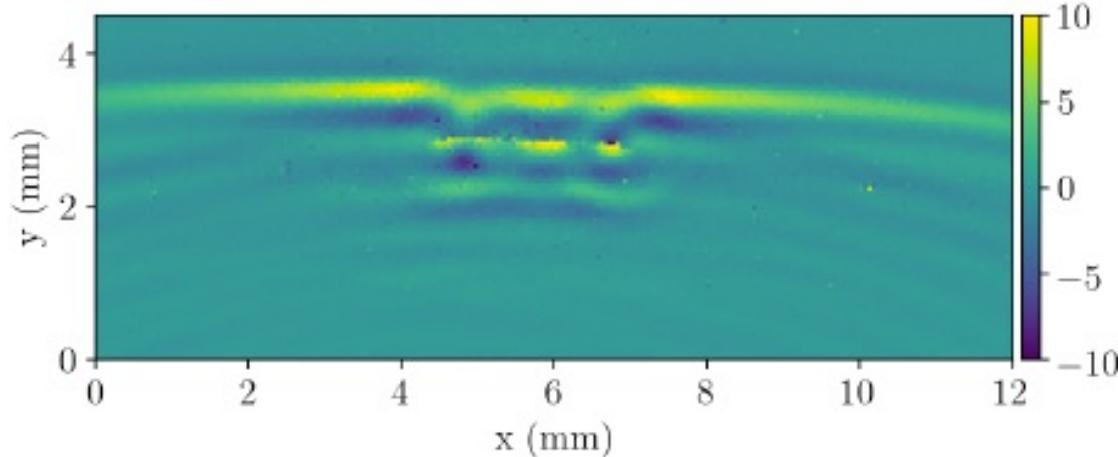
Where  $\phi$  are unknown, underlying PDE parameters we wish to invert for, and  $\{\underline{x}_l, \underline{u}_l\}$  are a set of (potentially noisy) observational data



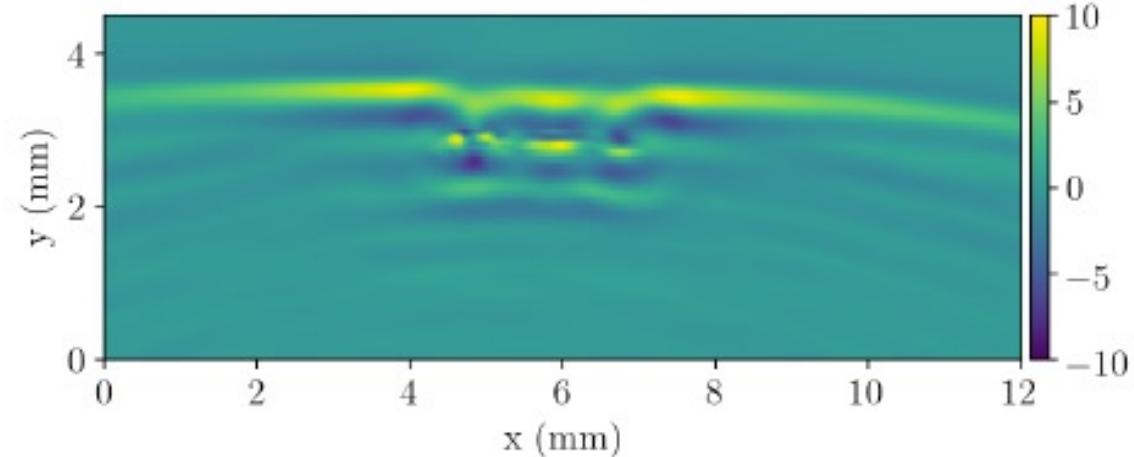
We **simultaneously** learn  $\theta$  and  $\phi$  when training the PINN

# PINNs for wave equation

Shukla et al, Physics-Informed Neural Network for Ultrasound Nondestructive Quantification of Surface Breaking Cracks, Journal of Nondestructive Evaluation (2020)



(a) Actual data at  $t = 12.38 \mu\text{s}$ .

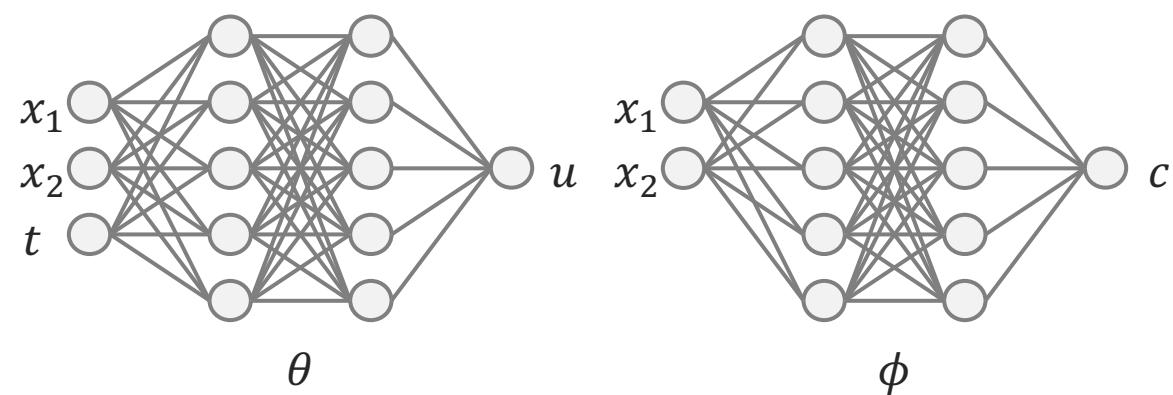


(b) Data recovered from PINN simulation at  $t = 12.38 \mu\text{s}$ .

$$L_p(\theta, \phi) = \frac{1}{N_p} \sum_i^{N_p} \left( \left[ \nabla^2 - \frac{1}{c(x_i; \phi)^2} \frac{\partial^2}{\partial t^2} \right] NN(x_i, t_i; \theta) \right)^2$$

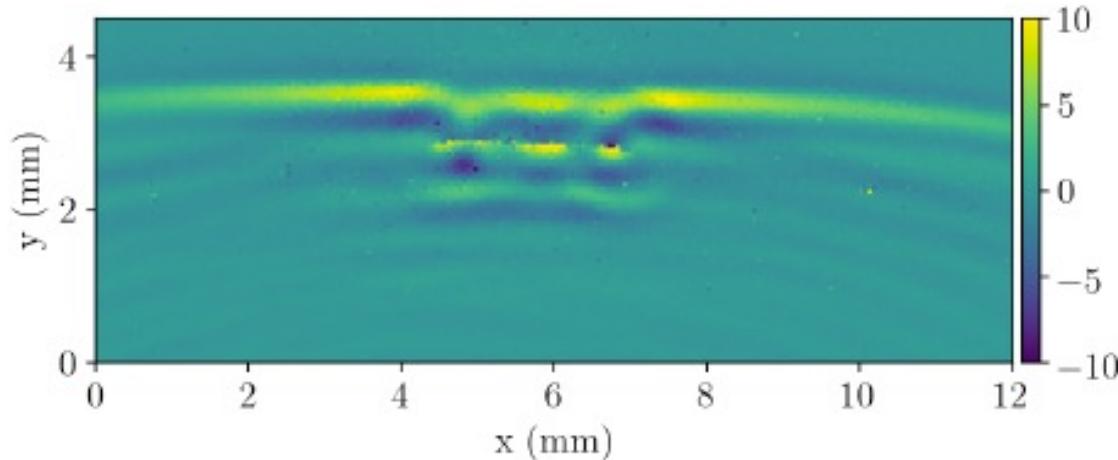
$$L_d(\theta) = \frac{\lambda}{N_d} \sum_l^{N_d} (NN(x_l, t_l; \theta) - u_{obs\ l})^2$$

Treat velocity model as **another** neural network, and simultaneously learn it

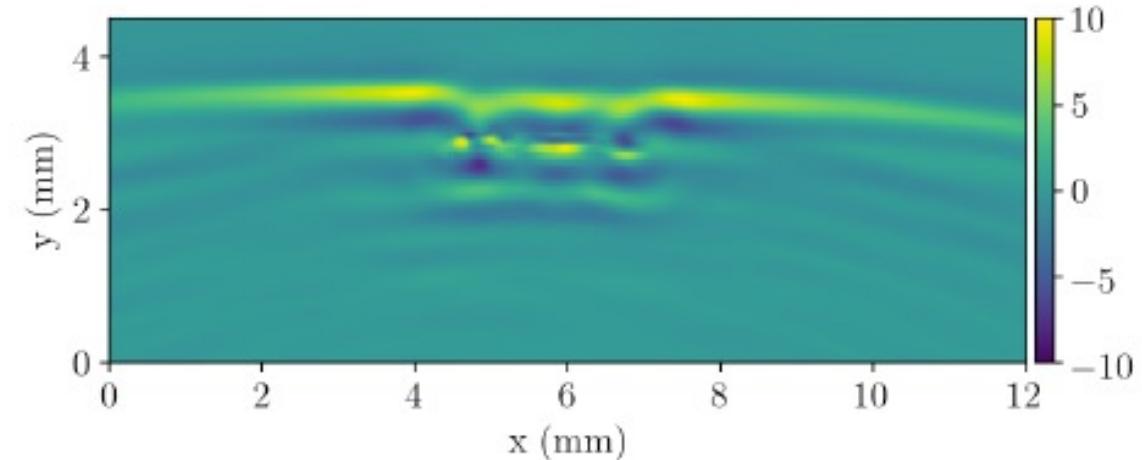


# PINNs for wave equation

Shukla et al, Physics-Informed Neural Network for Ultrasound Nondestructive Quantification of Surface Breaking Cracks, Journal of Nondestructive Evaluation (2020)



(a) Actual data at  $t = 12.38 \mu s$ .

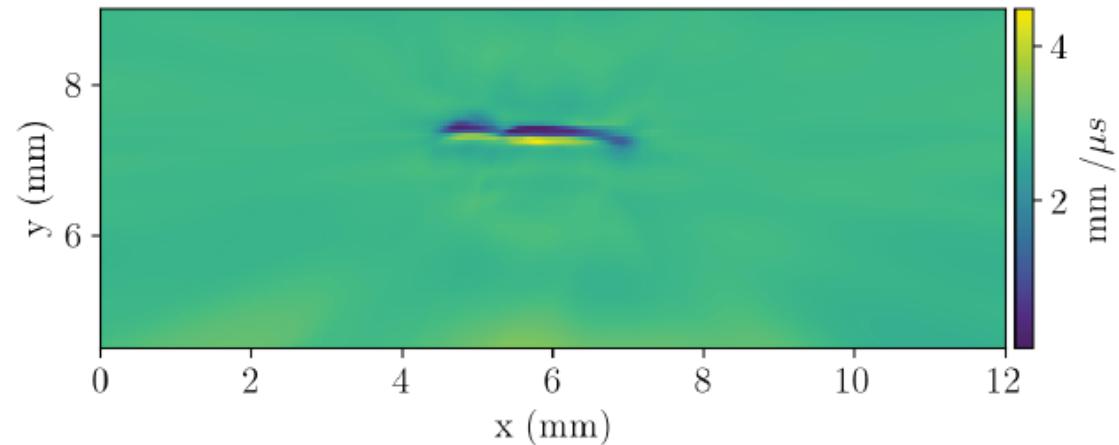


(b) Data recovered from PINN simulation at  $t = 12.38 \mu s$ .

$$L_p(\theta, \phi) = \frac{1}{N_p} \sum_i^{N_p} \left( \left[ \nabla^2 - \frac{1}{c(x_i; \phi)^2} \frac{\partial^2}{\partial t^2} \right] NN(x_i, t_i; \theta) \right)^2$$

$$L_d(\theta) = \frac{\lambda}{N_d} \sum_l^{N_d} (NN(x_l, t_l; \theta) - u_{obs\ l})^2$$

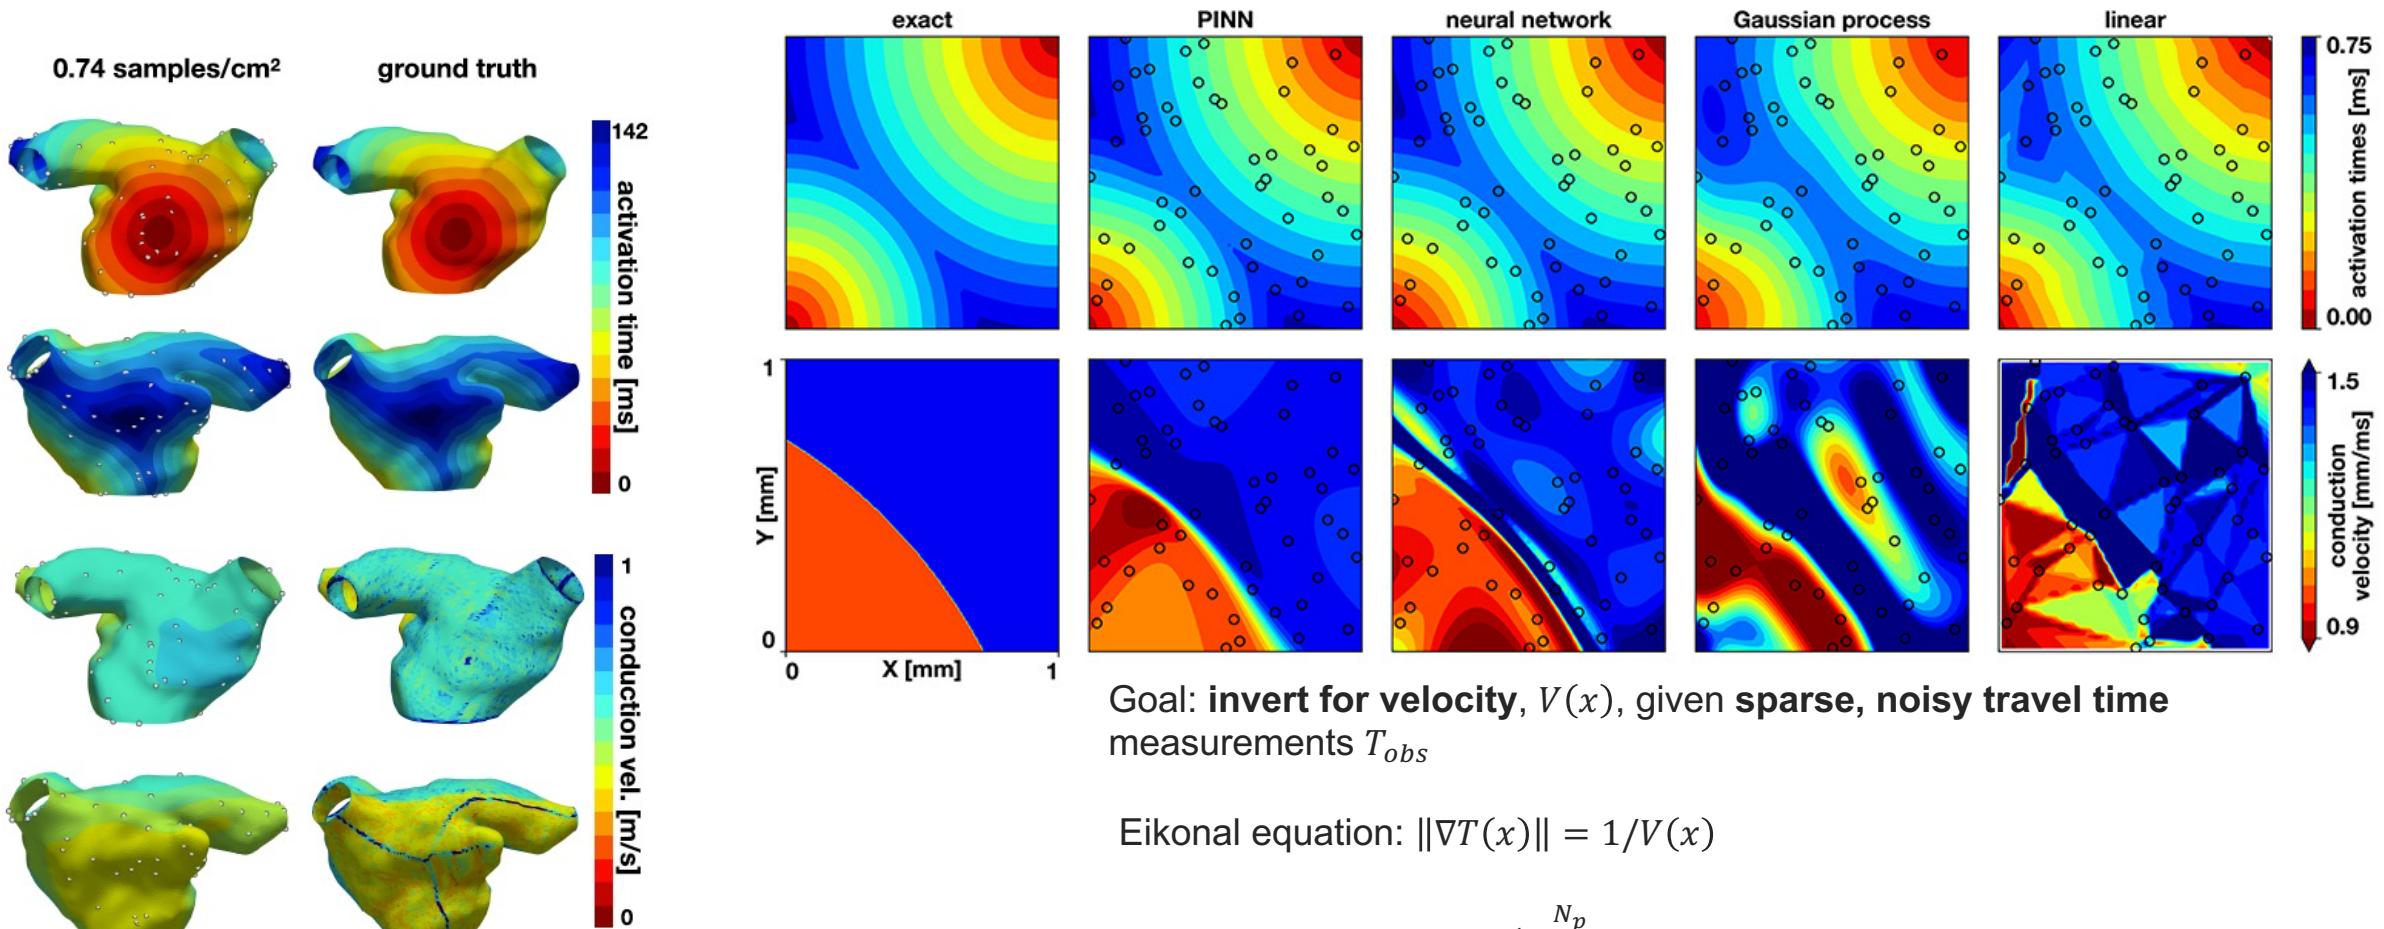
Treat velocity model as **another** neural network, and simultaneously learn it



(d) Speed  $v(x, y)$  recovered from PINN simulation.

# Live-coding a PINN in PyTorch

# PINNs for cardiac activation mapping



Goal: invert for velocity,  $V(x)$ , given sparse, noisy travel time measurements  $T_{obs}$

Eikonal equation:  $\|\nabla T(x)\| = 1/V(x)$

$$L_p(\theta, \phi) = \frac{1}{N_p} \sum_i^{N_p} (\|\nabla NN(x_i; \theta)\| V(x_i; \phi) - 1)^2$$

$$L_d(\theta) = \frac{\lambda}{N_d} \sum_l^{N_d} (NN(x_l; \theta) - T_{obs,l})^2$$

Sahli Costabal et al, Physics-Informed Neural Networks for Cardiac Activation Mapping, Frontiers in Physics (2020)

# PINNs for equation discovery

PINNs for solving **inverse** problems:

$$L(\theta, \phi) = L_p(\theta, \phi) + L_d(\theta)$$

$$L_p(\theta, \phi) = \frac{1}{N_p} \sum_i^{N_p} \|\mathcal{D}[NN(x_i; \theta); \phi] - f(x_i)\|^2 \quad \text{Physics loss}$$

$$L_d(\theta) = \frac{\lambda}{N_d} \sum_l^{N_d} \|NN(\mathbf{x}_l; \theta) - \mathbf{u}_l\|^2 \quad \text{Data loss}$$

Where  $\{\mathbf{x}_l, \mathbf{u}_l\}$  are a set of (potentially noisy) observational data

But how do we learn the **entire** differential operator  $\mathcal{D}$ , rather than its parameters  $\phi$ ?

# PINNs for equation discovery

How do we learn an **entire** differential operator  $\mathcal{D}$ ?

Build a **library** of  $n$  operators, such as:

$$\phi = (1, \partial_x, \partial_t, \partial_{xx}, \partial_{tt}, \partial_{xt})^T$$

Then assume the differential operator can be represented as

$$\mathcal{D} = \Lambda \phi$$

Where  $\Lambda$  is a (sparse) matrix of shape  $(d_u, n)$

E.g. for 1D damped harmonic oscillator:

$$\begin{aligned}\mathcal{D} &= (\mathbf{k} \quad \boldsymbol{\mu} \quad \mathbf{m} \quad \mathbf{0}) \begin{pmatrix} 1 \\ d_t \\ d_{tt} \\ d_{ttt} \end{pmatrix} \\ &= \mathbf{m} \frac{d^2}{dt^2} + \boldsymbol{\mu} \frac{d}{dt} + \mathbf{k}\end{aligned}$$

# PINNs for equation discovery

How do we learn an **entire** differential operator  $\mathcal{D}$ ?

Build a **library** of  $n$  operators, such as:

$$\phi = (1, \partial_x, \partial_t, \partial_{xx}, \partial_{tt}, \partial_{xt})^T$$

Then assume the differential operator can be represented as

$$\mathcal{D} = \Lambda \phi$$

Where  $\Lambda$  is a (sparse) matrix of shape  $(d_u, n)$

E.g. for 1D damped harmonic oscillator:

$$\begin{aligned}\mathcal{D} &= (\mathbf{k} \quad \boldsymbol{\mu} \quad \mathbf{m} \quad \mathbf{0}) \begin{pmatrix} 1 \\ d_t \\ d_{tt} \\ d_{ttt} \end{pmatrix} \\ &= \mathbf{m} \frac{d^2}{dt^2} + \boldsymbol{\mu} \frac{d}{dt} + \mathbf{k}\end{aligned}$$

PINNs for **equation discovery**:

$$L(\theta, \Lambda) = L_p(\theta, \Lambda) + L_d(\theta)$$

$$L_p(\theta, \Lambda) = \frac{1}{N_p} \sum_i^{N_p} \|\Lambda \phi[NN(x_i; \theta)]\|^2 + \|\Lambda\|^2 \quad \text{Physics loss}$$

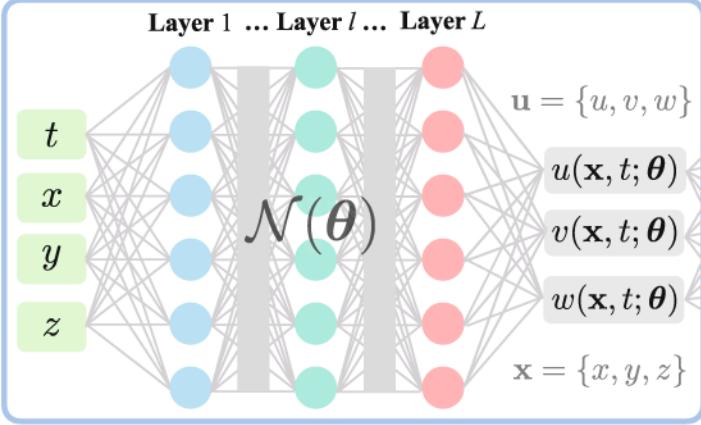
$$L_d(\theta) = \frac{\lambda}{N_d} \sum_l^{N_d} \|NN(\mathbf{x}_l; \theta) - \mathbf{u}_l\|^2 \quad \text{Data loss}$$

Where  $\Lambda$  are treated as **learnable** parameters and  $\{\mathbf{x}_l, \mathbf{u}_l\}$  are a set of (potentially noisy) observational data

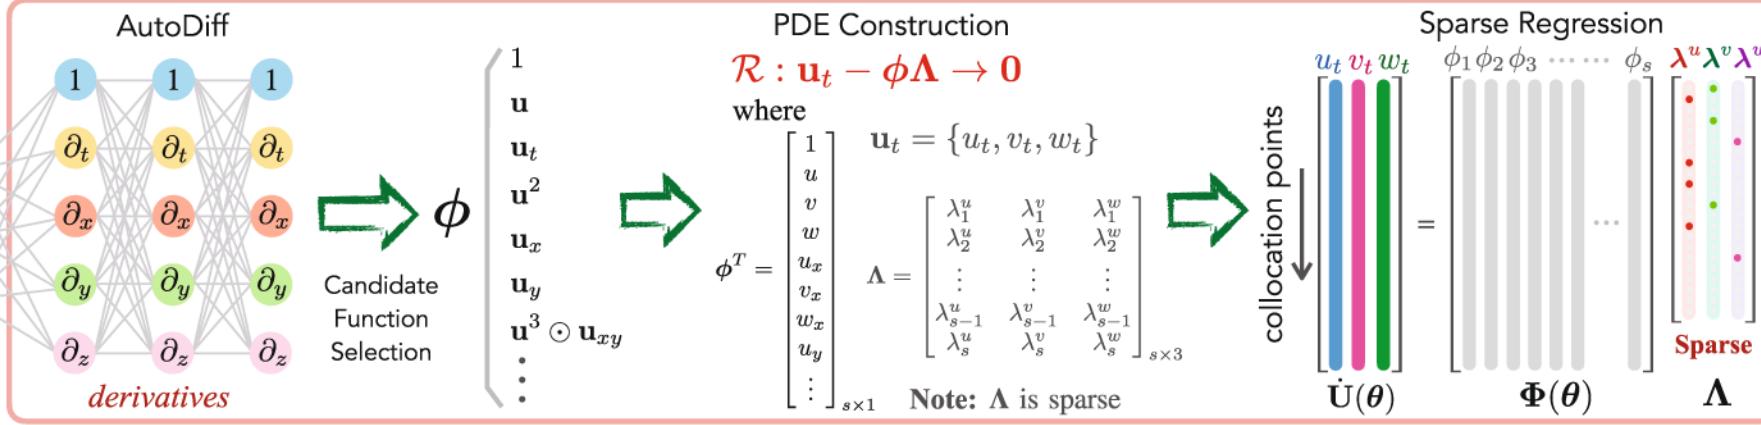
Typically, some regularization / prior on  $\Lambda$  (e.g. sparsity) is needed, as this optimisation problem can be very **ill-posed**

# PINNs for equation discovery

DNN with Unknown Parameters  $\theta$



Physical Law with Unknown Parameters  $\Lambda$



$$\text{Data Loss: } \underbrace{\mathcal{L}_d(\theta; \mathcal{D}_u)}_{\text{measurement}} = \frac{1}{N_m} \|\mathbf{u}^\theta - \mathbf{u}^m\|_2^2 \rightarrow \underbrace{\mathcal{L}(\theta, \Lambda; \mathcal{D}_u, \mathcal{D}_c)}_{\text{total loss}} = \underbrace{\mathcal{L}_d(\theta; \mathcal{D}_u)}_{\text{data loss}} + \underbrace{\alpha \mathcal{L}_p(\theta, \Lambda; \mathcal{D}_c)}_{\text{physics loss}} + \underbrace{\beta \|\Lambda\|_0}_{\text{regularization}} \leftarrow \text{Residual Loss: } \mathcal{L}_p(\theta, \Lambda; \mathcal{D}_c) = \frac{1}{N_c} \|\mathbf{U}(\theta) - \Phi(\theta) \Lambda\|$$

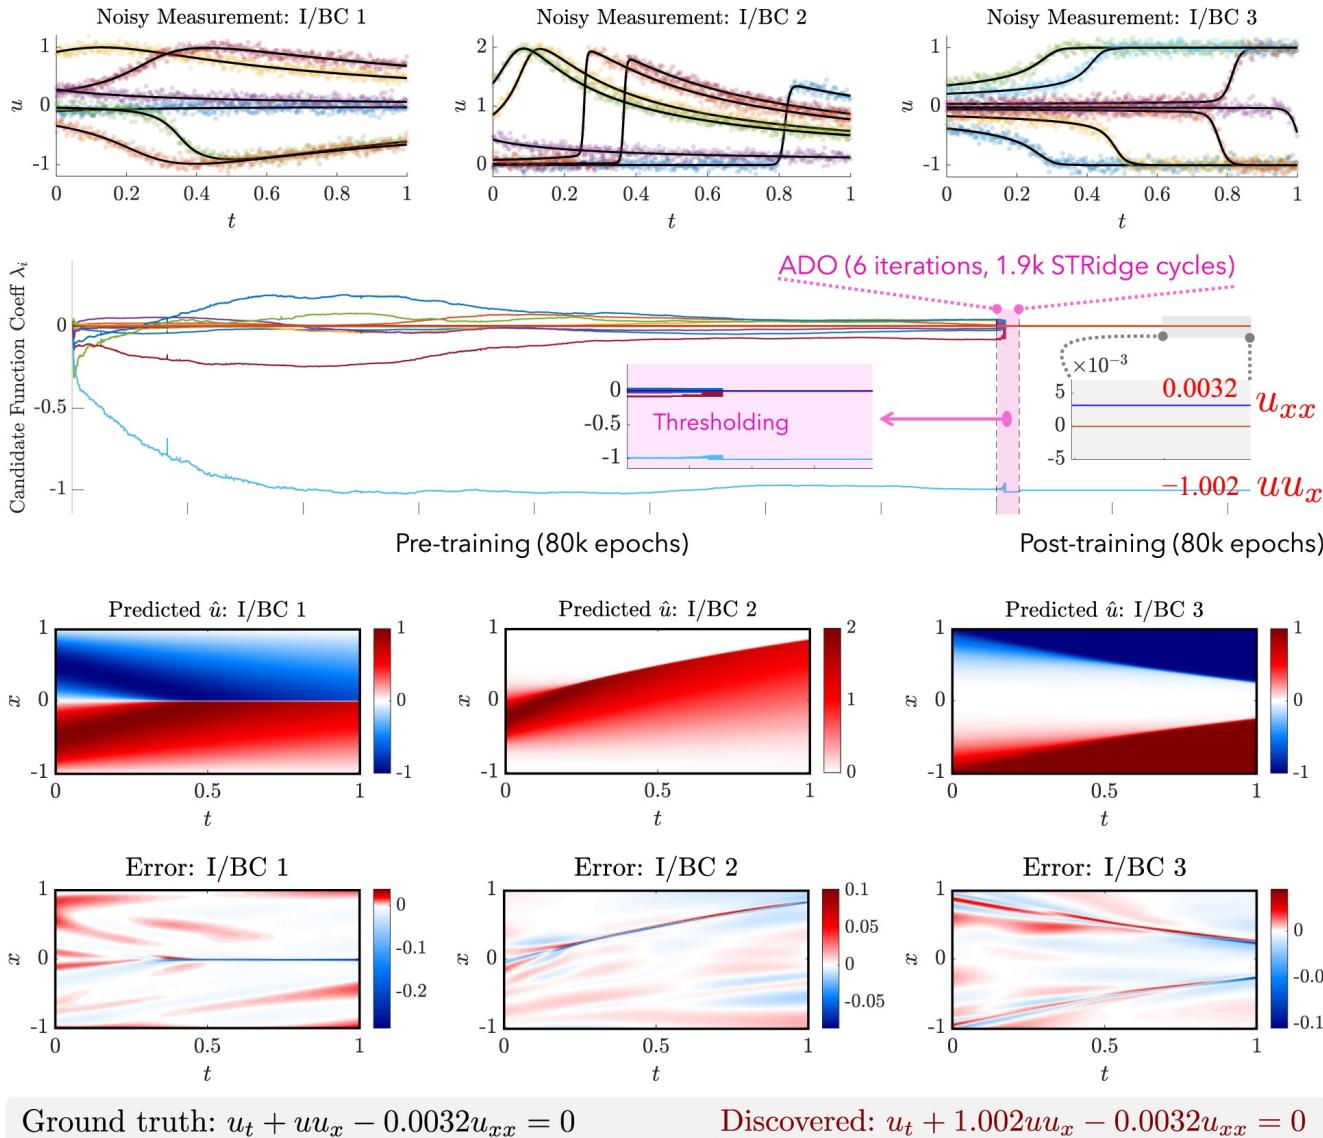
$$\text{Solution by ADO: } \hat{\Lambda}_{k+1} := \arg \min_{\Lambda} \left[ \|\dot{\mathbf{U}}(\hat{\theta}_k) - \Phi(\hat{\theta}_k) \Lambda\|_2^2 + \beta \|\Lambda\|_0 \right] \text{ by STRidge} \quad \hat{\theta}_{k+1} := \arg \min_{\theta} [\mathcal{L}_d(\theta; \mathcal{D}_u) + \alpha \mathcal{L}_p(\theta, \hat{\Lambda}_{k+1}; \mathcal{D}_c)] \text{ by DNN training}$$



- Trains by alternating between updating  $\Lambda$  and  $\theta$

Chen et al, Physics-informed learning of governing equations from scarce data, Nature communications (2021)

# PINNs for equation discovery



- PINN “discovering” Burgers’ equation
- By combining datasets sampled under three different I/BCs with 10% noise

Chen et al, Physics-informed learning of governing equations from scarce data, Nature communications (2021)