



Deep Learning in Scientific Computing

Introduction to Differentiable Physics - Part 2

Spring Semester 2023

Siddhartha Mishra
Ben Moseley

ETH zürich

Course timeline

Tutorials

Tue 3:15-14:00, HG E5

21.02.

28.02. ~~Intro to PyTorch~~

07.03. ~~Deep learning in PyTorch I~~

14.03. ~~Deep learning in PyTorch II~~

21.03. ~~Implementing PINNs I~~

28.03. ~~Implementing PINNs II~~

04.04. ~~Implementing PINNs III~~

11.04.

18.04. ~~Introduction to projects~~

25.04. ~~Implementing neural operators I~~

02.05. ~~Implementing neural operators II~~

09.05. ~~Project work~~

16.05. ~~Implementing neural operators III~~

23.05. ~~Project work~~

30.05. ~~Coding an autodiff engine~~

Lectures

Fri 12:15-14:00, HG D1.1

24.02. ~~Course introduction~~

03.03. ~~Introduction to deep learning I~~

10.03. ~~Introduction to deep learning II~~

17.03. ~~Physics-informed neural networks—introduction and theory~~

24.03. ~~Physics-informed neural networks—applications~~

31.03. ~~Physics-informed neural networks—limitations and extensions~~

07.04.

14.04.

21.04. ~~Introduction to operator learning~~

28.04. ~~Operator networks and DeepONet~~

05.05. ~~DeepONet continuation~~

12.05. ~~Neural operators~~

19.05. ~~Limitations of neural operators~~

26.05. ~~Introduction to differentiable physics I~~

02.06. ~~Introduction to differentiable physics II~~

Lecture overview

- Differentiable physics recap
- Coding a simple hybrid approach in PyTorch
- Hybrid approaches for inverse problems
- Neural differential equations (NDEs)
- Course summary

Hybrid approaches - recap

Advantages of DNNs

- Usually very **fast** (once trained)
- Can represent highly **non-linear** functions

Limitations of DNNs

- Often lots of **training data required**
- Can be hard to **optimise**
- Can be hard to **interpret**
- Often struggle to **generalise**

General advice

Use DNNs to:

- 1) **Accelerate** your workflow, or
- 2) Learn the **parts** you are unsure of / have incomplete knowledge

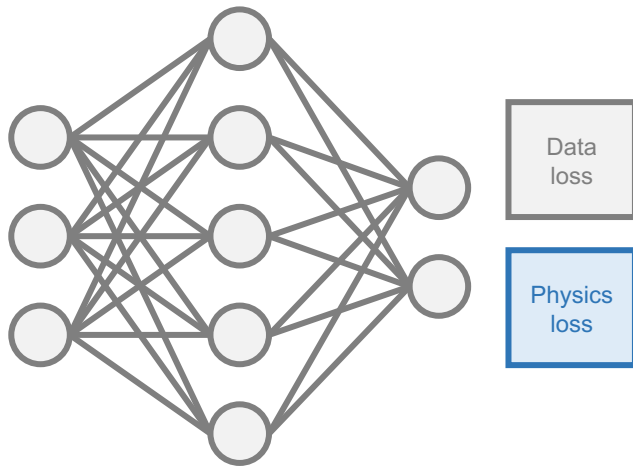
Otherwise using DNNs may **not** be a good idea!



Key idea: incorporate DNNs directly into a traditional algorithm
= **hybrid approach**

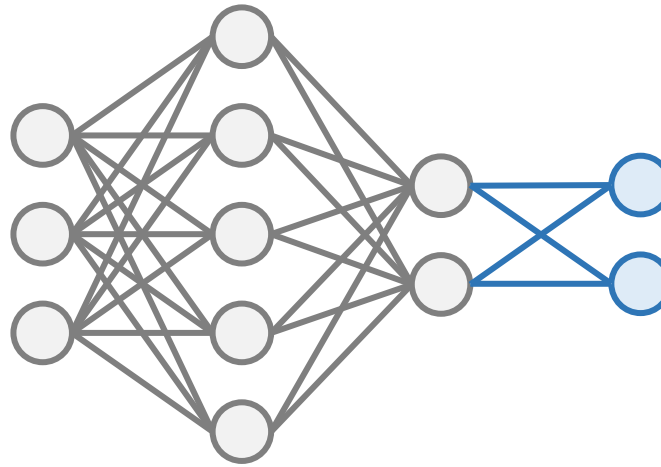
Ways to incorporate scientific principles into machine learning

Loss function



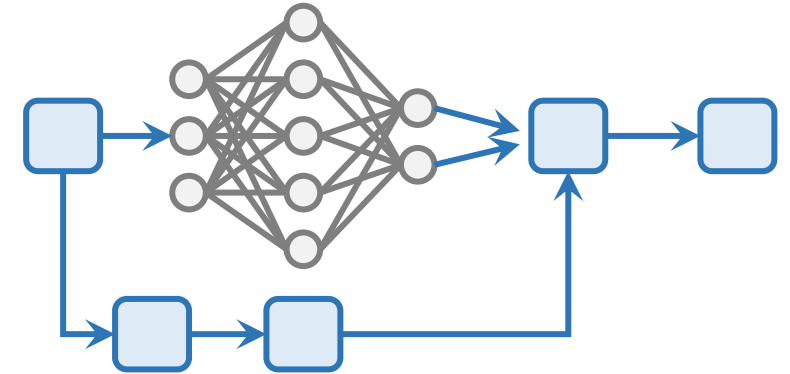
Example:
Physics-informed neural networks
(add governing equations to loss function)

Architecture



Example:
Encoding regularity / symmetries /
conservation laws (e.g. energy conservation,
rotational invariance), **operator learning**

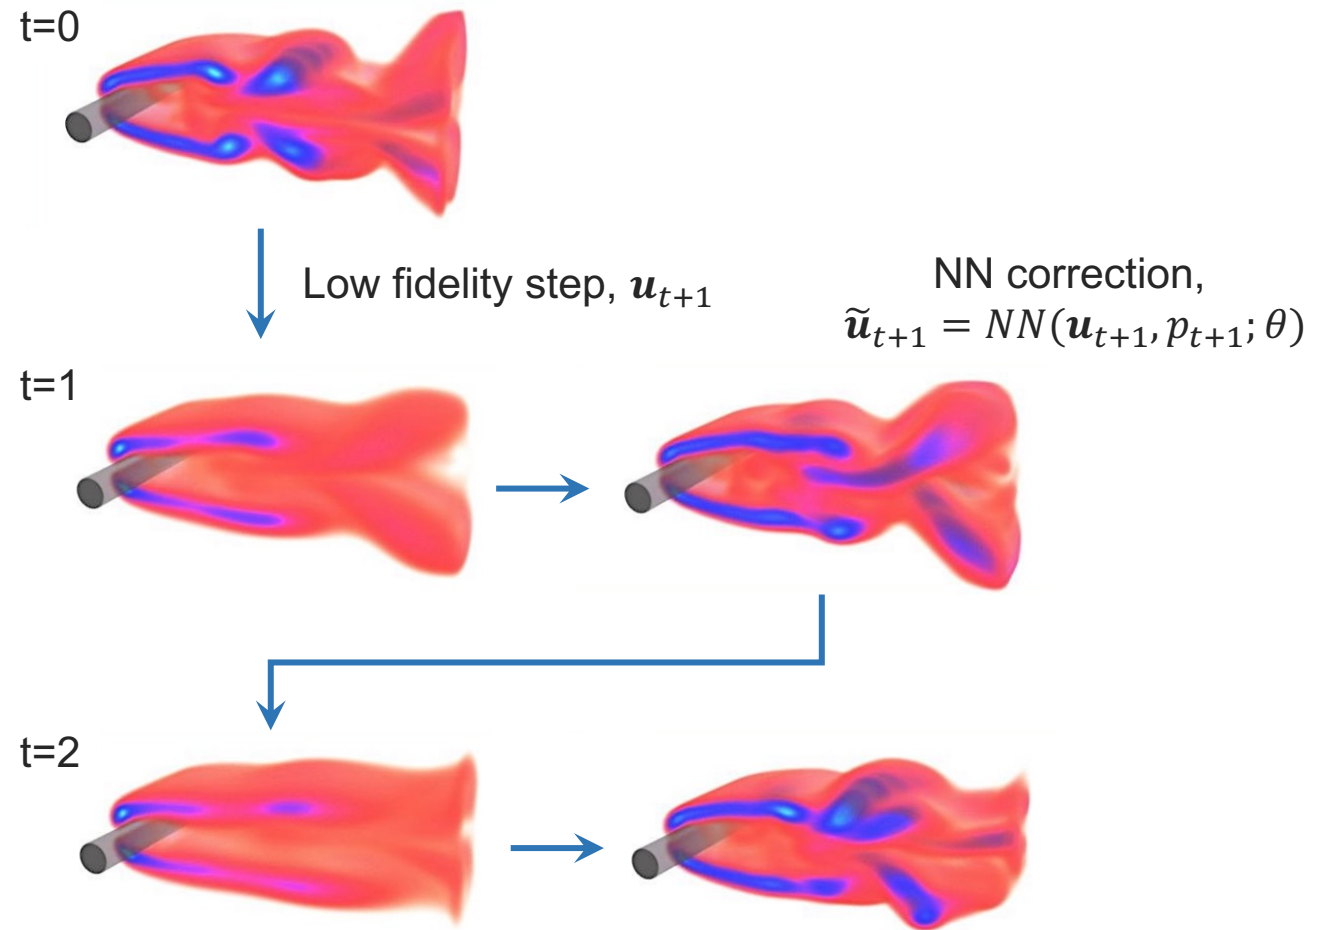
Hybrid approaches



Example:
Neural differential equations
(incorporating neural networks into
traditional PDE solvers)

Hybrid Navier-Stokes solver

```
def NS_solver(u_0, p_0, rho, nu):  
    "Pseudocode for solving NS equation"  
  
    # u_0, p_0 have shape (NX, NY, NZ)  
    u_t, p_t = u_0, p_0  
    for t in range(0, T):  
        u_star = f(u_t, p_t, rho, nu)  
        p_t = matrix_solve(u_star, p_t, rho)  
        u_t = g(u_t, p_t, rho, nu)  
  
    return u_t, p_t  
  
def Hybrid_NS_solver(u_0, p_0, rho, nu, theta):  
    "Pseudocode for solving NS equation, with NN correction"  
  
    # u_0, p_0 have shape (NX, NY, NZ)  
    u_t, p_t = u_0, p_0  
    for t in range(0, T):  
        u_star = f(u_t, p_t, rho, nu)  
        p_t = matrix_solve(u_star, p_t, rho)  
        u_t = g(u_t, p_t, rho, nu)  
  
        u_t, p_t = NN(u_t, p_t, theta)  
  
    return u_t, p_t
```



Um et al, Solver-in-the-loop: Learning from differentiable physics to interact with iterative PDE-solvers, NeurIPS (2020)

How do we train hybrid approaches?



Key idea: **autodifferentiation** allows us to differentiate **arbitrary** algorithms, not just neural networks!

Differentiable physics = using autodifferentiation to learn physical algorithms


```
def NN(x, theta):  
    "Defines a FCN"  
    y = torch.tanh(theta[0]*x + theta[1])  
    return y
```

```
theta.requires_grad_(True)  
y = NN(x, theta)  
loss = loss_fn(y, y_true)  
dtheta = torch.autograd(loss, theta)  
# for learning theta (training NN)
```

```
def Hybrid_NS_solver(u_0, p_0, rho, nu, theta):  
    "Pseudocode for solving NS equation, with NN correction"  
  
    # u_0, p_0 have shape (NX, NY, NZ)  
    u_t, p_t = u_0, p_0  
    for t in range(0, T):  
        u_star = f(u_t, p_t, rho, nu)  
        p_t = matrix_solve(u_star, p_t, rho)  
        u_t = g(u_t, p_t, rho, nu)  
  
        u_t, p_t = NN(u_t, p_t, theta)  
  
    return u_t, p_t
```

```
theta.requires_grad_(True)  
u_T, _ = Hybrid_NS_solver(u_0, p_0, rho, nu, theta)  
loss = loss_fn(u_T, u_T_true)  
dtheta = torch.autograd(loss, theta)  
# for learning theta (training NN)
```

Hybrid approach - results



Um et al, Solver-in-the-loop: Learning from differentiable physics to interact with iterative PDE-solvers, NeurIPS (2020)

Lecture overview

- Differentiable physics recap
- Coding a simple hybrid approach in PyTorch
- Hybrid approaches for inverse problems
- Neural differential equations (NDEs)
- Course summary

Coding a simple hybrid approach in PyTorch

More complex hybrid approaches

```
def NS_solver(u_0, p_0, rho, nu):
    "Pseudocode for solving NS equation"

    # u_0, p_0 have shape (NX, NY, NZ)
    u_t, p_t = u_0, p_0
    for t in range(0, T):
        u_star = f(u_t, p_t, rho, nu)
        p_t = matrix_solve(u_star, p_t, rho)
        u_t = g(u_t, p_t, rho, nu)

    return u_t, p_t

def Hybrid_NS_solver(u_0, p_0, rho, nu, theta):
    "Pseudocode for solving NS equation, with NN correction"

    # u_0, p_0 have shape (NX, NY, NZ)
    u_t, p_t = u_0, p_0
    for t in range(0, T):
        u_star = f(u_t, p_t, rho, nu)
        p_t = matrix_solve(u_star, p_t, rho)
        u_t = g(u_t, p_t, rho, nu)

        u_t, p_t = NN(u_t, p_t, theta)

    return u_t, p_t
```

- There are many **other** ways we can insert neural networks into our existing algorithms

More complex hybrid approaches

```
def NS_solver(u_0, p_0, rho, nu):
    "Pseudocode for solving NS equation"

    # u_0, p_0 have shape (NX, NY, NZ)
    u_t, p_t = u_0, p_0
    for t in range(0, T):
        u_star = f(u_t, p_t, rho, nu)
        p_t = matrix_solve(u_star, p_t, rho)
        u_t = g(u_t, p_t, rho, nu)

    return u_t, p_t

def Hybrid_NS_solver(u_0, p_0, rho, nu, theta):
    "Pseudocode for solving NS equation, with NN correction"

    # u_0, p_0 have shape (NX, NY, NZ)
    u_t, p_t = u_0, p_0
    for t in range(0, T):
        u_star = f(u_t, p_t, rho, nu)
        p_t = matrix_solve(u_star, p_t, rho)
        u_t = g(u_t, p_t, rho, nu)


        u_t, p_t = NN(u_t, p_t, theta)

    return u_t, p_t
```


- There are many **other** ways we can insert neural networks into our existing algorithms
- Many traditional algorithms (simulation, inversion, control, data assimilation, equation discovery, ...) are **iterative**
- “**In-the-loop**” methods are a class of hybrid approaches which insert neural networks into the **inner loops** of traditional algorithms

A more complex “in-the-loop” example


Consider the following **inverse** problem:



Ground truth computed tomography image



Resulting tomographic data (sinogram)



Result of inverse algorithm





Image source: Wikipedia

$$a(x) \quad b = F(a) = I_0 \exp\left(-\int_l a(x) dx\right) \quad \hat{a}(x)$$



$$b = F(a)$$

a = set of input conditions


F = physical model of the system

b = resulting properties given F and a


Adler et al, Solving ill-posed inverse problems using iterative deep neural networks, Inverse Problems (2017)

A more complex “in-the-loop” example


Consider the following **inverse** problem:



Ground truth computed tomography image




Resulting tomographic data (sinogram)



Result of inverse algorithm

$$a(x) \qquad b = F(a) = I_0 \exp\left(-\int_l a(x) dx\right) \qquad \hat{a}(x)$$



Adler et al, Solving ill-posed inverse problems using iterative deep neural networks, Inverse Problems (2017)

This problem can be framed as an **optimisation** problem:

$$\min_{\hat{a}} \|b - F(\hat{a})\|^2$$

Assuming F is differentiable, we can use **gradient descent** to learn \hat{a} :

Loss function:

$$L(\hat{a}) = \|b - F(\hat{a})\|^2$$

Gradient descent:


$$\hat{a} \leftarrow \hat{a} - \gamma \frac{\partial L(\hat{a})}{\partial \hat{a}}$$

A more complex “in-the-loop” example


Consider the following **inverse** problem:

However, this problem can be very **ill-posed** if;


- The measurements are noisy
- There are not enough observations



Ground truth computed tomography image



Resulting tomographic data (sinogram)




Result of inverse algorithm

$a(x)$

$$b = F(a) = I_0 \exp\left(-\int_l a(x) dx\right)$$


$\hat{a}(x)$




Adler et al, Solving ill-posed inverse problems using iterative deep neural networks, Inverse Problems (2017)

A more complex “in-the-loop” example


Consider the following **inverse** problem:



Ground truth computed tomography image



Resulting tomographic data (sinogram)




Result of inverse algorithm

$a(x)$

$$b = F(a) = I_0 \exp\left(-\int_l a(x) dx\right)$$

$\hat{a}(x)$



Adler et al, Solving ill-posed inverse problems using iterative deep neural networks, Inverse Problems (2017)

However, this problem can be very **ill-posed** if;

- The measurements are noisy
- There are not enough observations

To improve, add **regularization**:

$$L(\hat{a}) = \|b - F(\hat{a})\|^2 + \lambda R(\hat{a})$$

Where, for example

$$R(\hat{a}) = \|\nabla \hat{a}\|$$

Which asserts a **prior** that the output image should be “smooth” (= total variation regularization)

A more complex “in-the-loop” example

```
def X_ray_tomography(a_hat_0, b):  
    "Pseudocode for carrying out X ray tomography"  
  
    # a_hat_0 is the initial image guess, of shape (NX, NY)  
    # b are the observed measurements, of shape (MX, MY)  
  
    a_hat = a_hat_0  
    lam = 1  
    for i in range(0, n_steps):  
        a_hat = a_hat.requires_grad_(True)  
        b_hat = numerical_integrate(a_hat)  
        R = total_variation(b_hat)  
        loss = torch.mean((b-b_hat)**2) + lam*R  
        da = torch.autograd.grad(loss, a_hat)  
        a_hat -= gamma*da  
  
    return a_hat
```

1. Start with initial guess \hat{a}
2. Loop:
 1. Compute gradient, $\frac{\partial L(\hat{a})}{\partial \hat{a}}$
 2. Take gradient descent step,

$$\hat{a} \leftarrow \hat{a} - \gamma \frac{\partial L(\hat{a})}{\partial \hat{a}}$$

However, this problem can be very **ill-posed** if;

- The measurements are noisy
- There are not enough observations

To improve, add **regularization**:

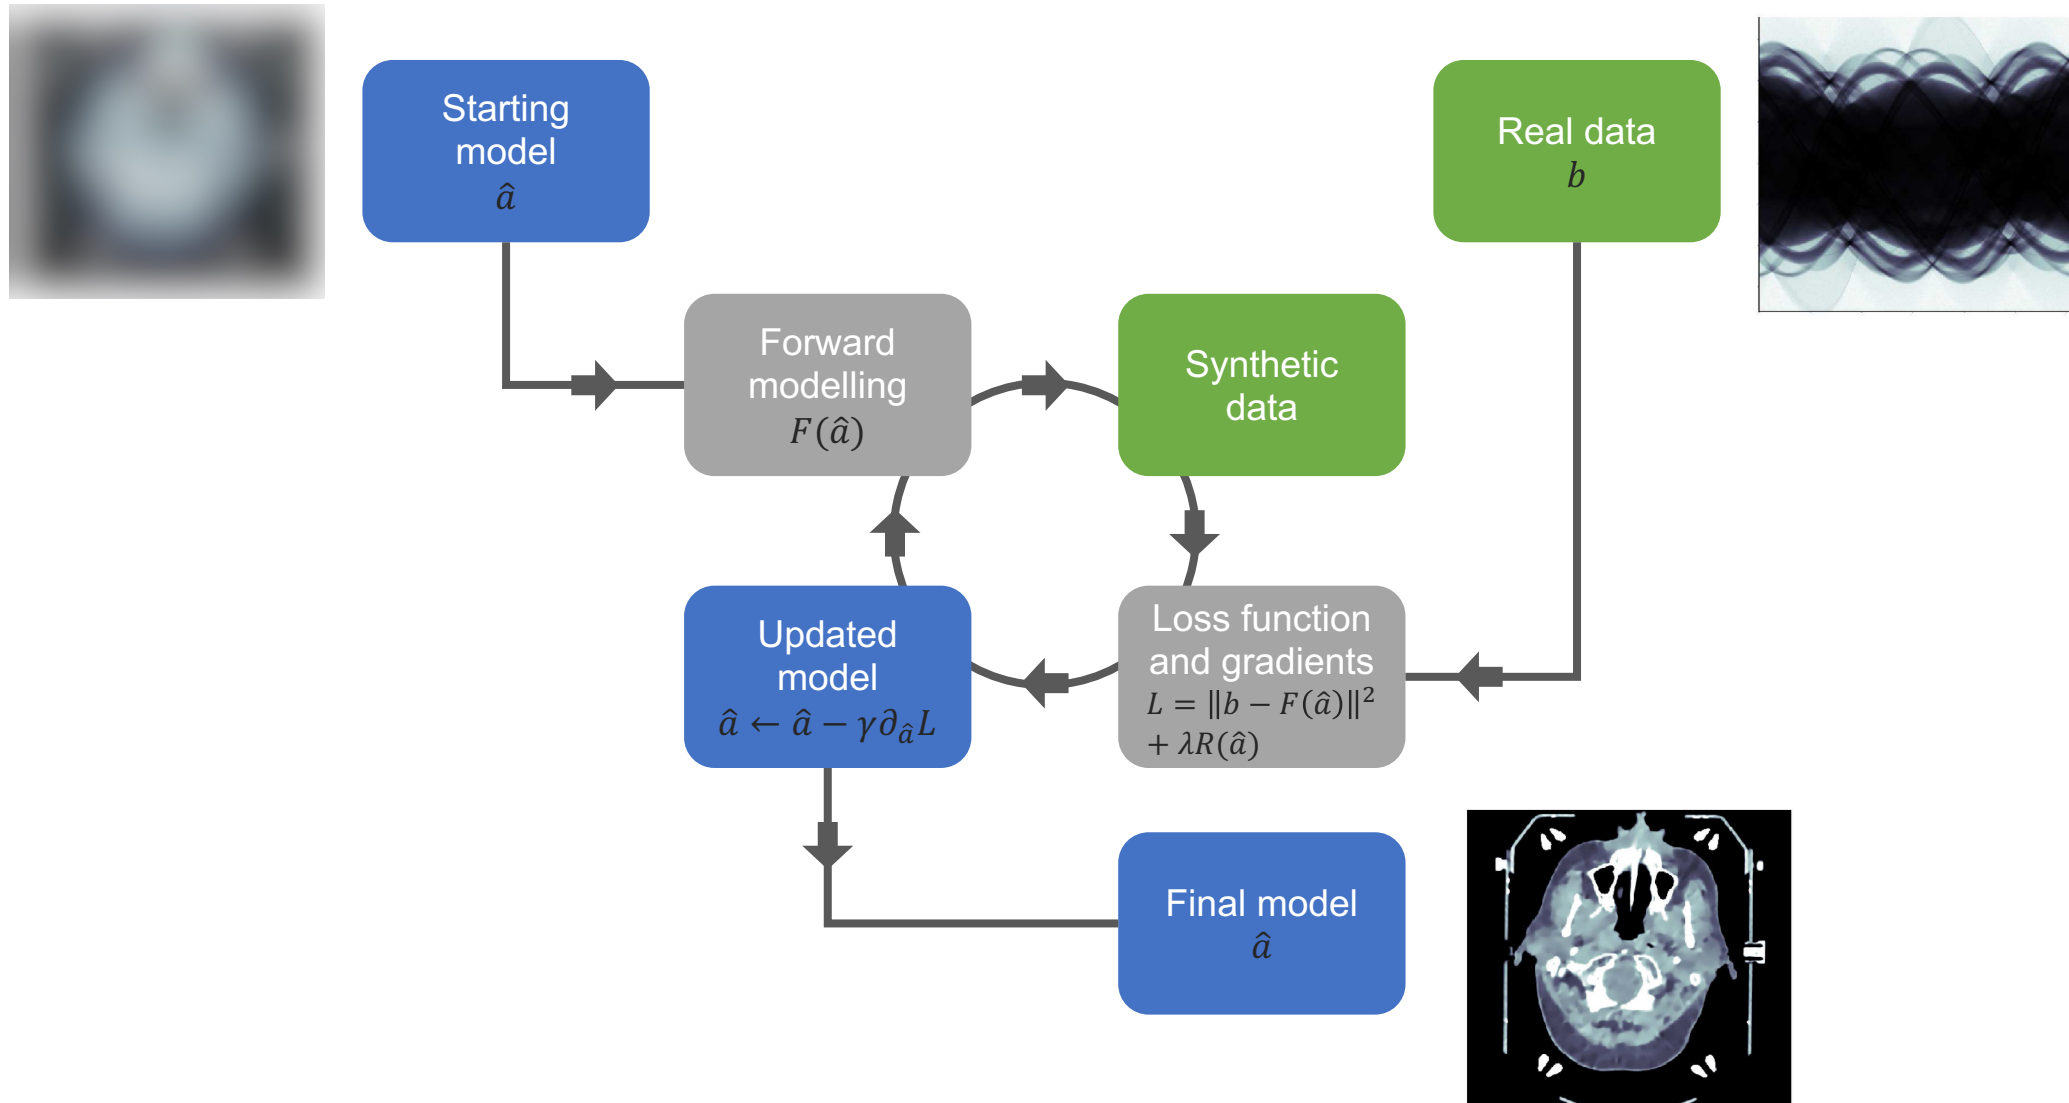
$$L(\hat{a}) = \|b - F(\hat{a})\|^2 + \lambda R(\hat{a})$$

Where, for example

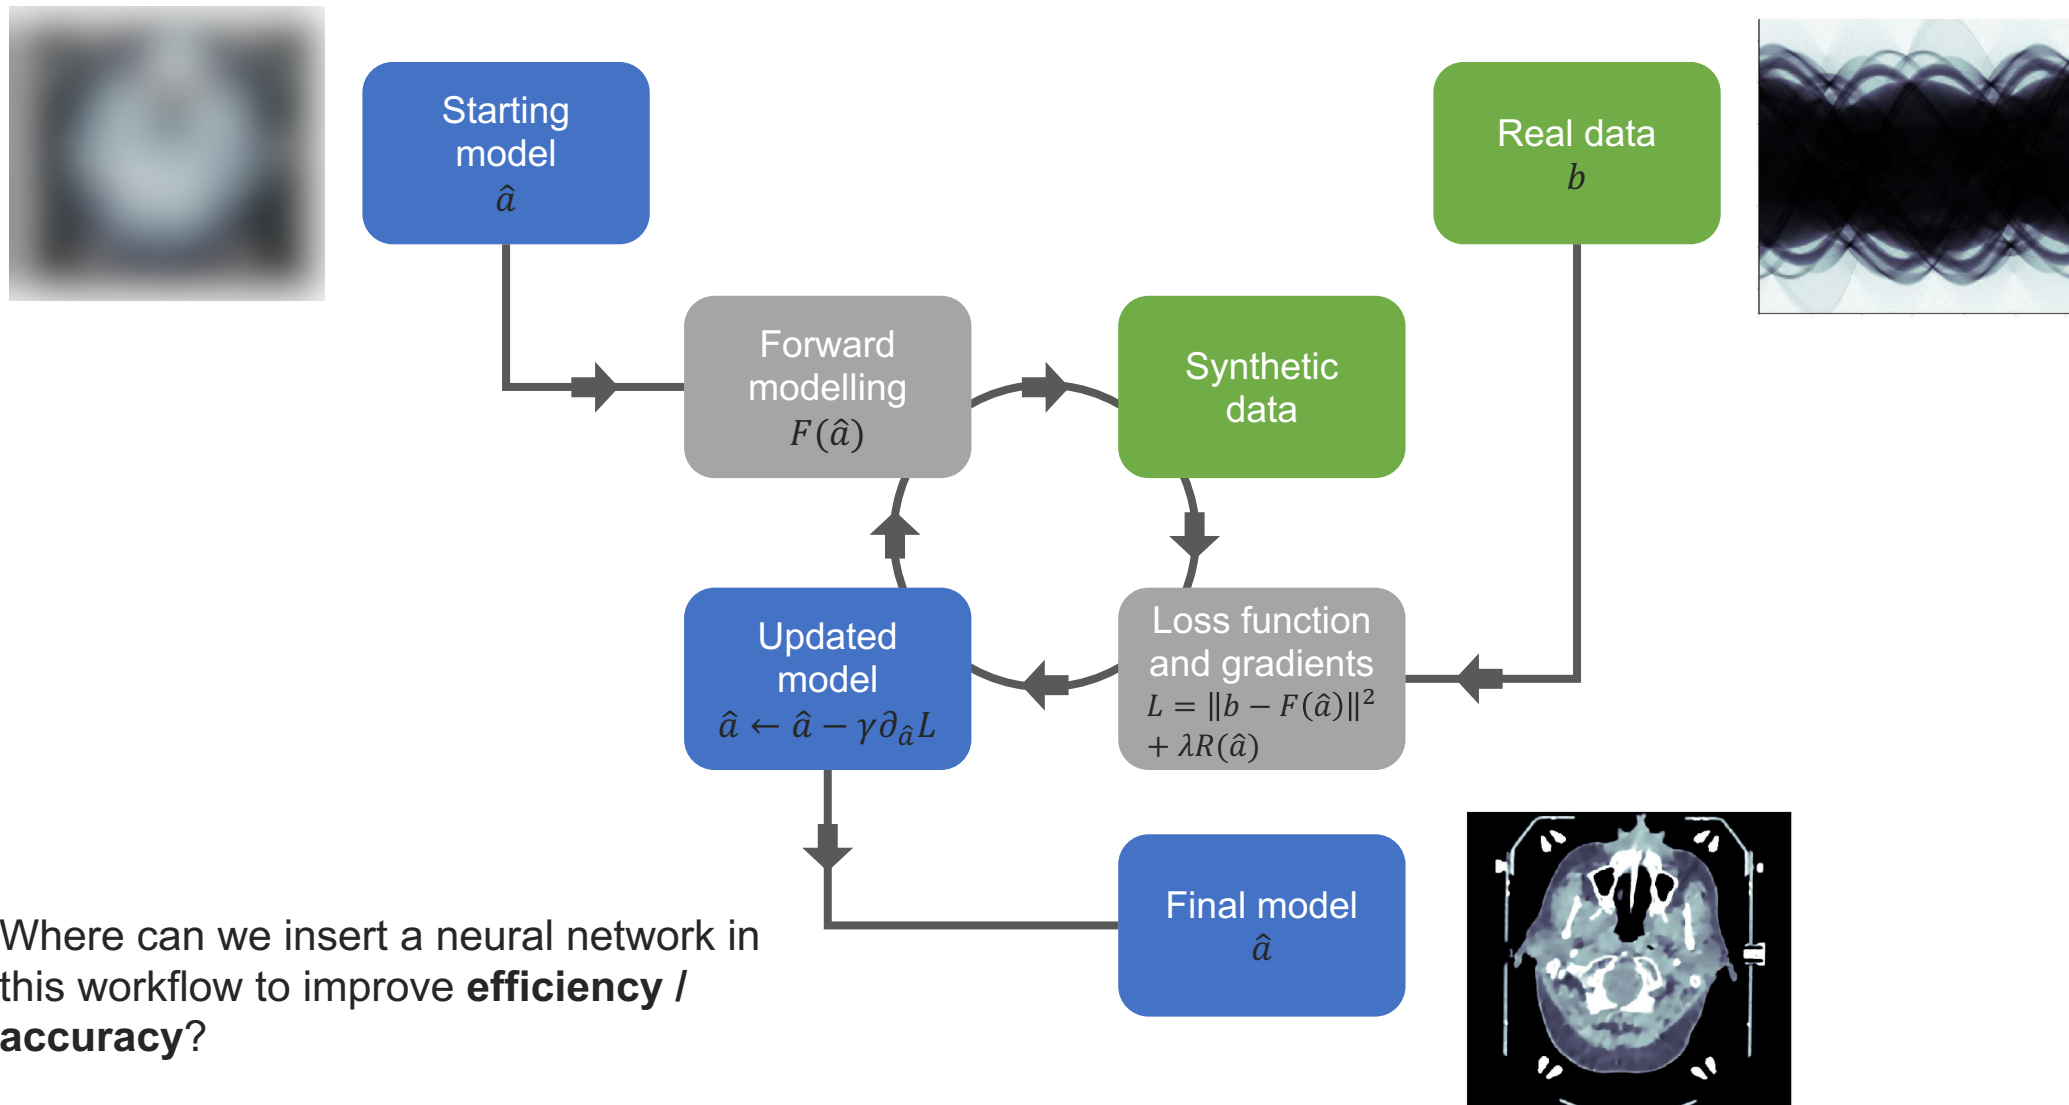
$$R(\hat{a}) = \|\nabla \hat{a}\|$$

Which asserts a **prior** that the output image should be “smooth” (= total variation regularization)

X-ray tomography

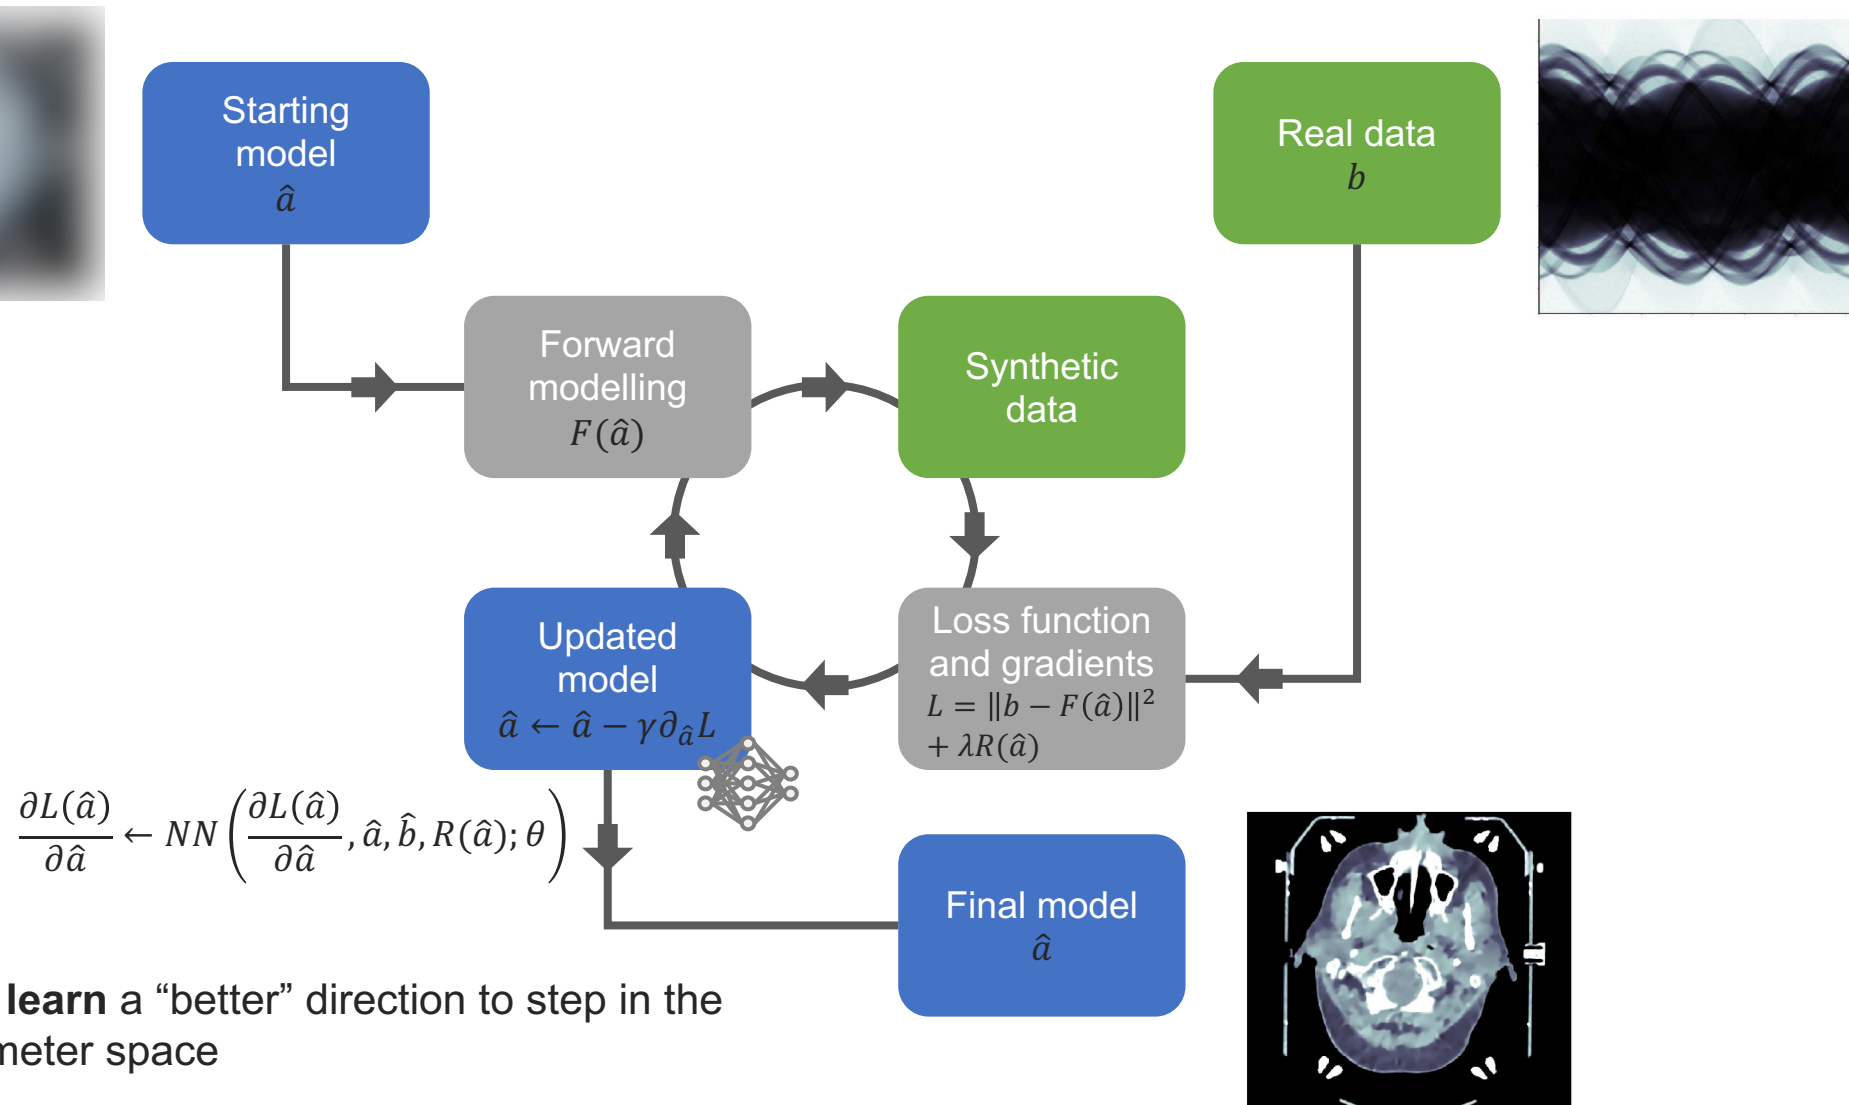


X-ray tomography



- Where can we insert a neural network in this workflow to improve **efficiency** / **accuracy**?

Hybrid X-ray tomography



Idea: **learn** a “better” direction to step in the parameter space

Hybrid X-ray tomography

```
def X_ray_tomography(a_hat_0, b):  
    "Pseudocode for carrying out X ray tomography"  
  
    # a_hat_0 is the initial image guess, of shape (NX, NY)  
    # b are the observed measurements, of shape (MX, MY)  
  
    a_hat = a_hat_0  
    lam = 1  
    for i in range(0, n_steps):  
        a_hat = a_hat.requires_grad_(True)  
        b_hat = numerical_integrate(a_hat)  
        R = total_variation(b_hat)  
        loss = torch.mean((b-b_hat)**2) + lam*R  
        da = torch.autograd.grad(loss, a_hat)  
        a_hat -= gamma*da  
  
    return a_hat
```

1. Start with initial guess \hat{a}

2. Loop:

1. Compute gradient, $\frac{\partial L(\hat{a})}{\partial \hat{a}}$

2. Take gradient descent step,

$$\hat{a} \leftarrow \hat{a} - \gamma \text{NN} \left(\frac{\partial L(\hat{a})}{\partial \hat{a}}, \hat{a}, \hat{b}, R(\hat{a}); \theta \right)$$

```
def Hybrid_X_ray_tomography(a_hat_0, b, theta):  
    "Pseudocode for carrying out X ray tomography, with NN correction"  
  
    # a_hat_0 is the initial image guess, of shape (NX, NY)  
    # b are the observed measurements, of shape (MX, MY)  
  
    a_hat = a_hat_0  
    lam = 1  
    for i in range(0, n_steps):  
        a_hat = a_hat.requires_grad_(True)  
        b_hat = numerical_integrate(a_hat)  
        R = total_variation(b_hat)  
        loss = torch.mean((b-b_hat)**2) + lam*R  
        da = torch.autograd.grad(loss, a_hat)  
        da = NN(da, a_hat, b_hat, R, theta)  
        a_hat -= gamma*da  
  
    return a_hat
```



Idea: **learn** a “better” direction to step in the parameter space

Hybrid X-ray tomography

```
def X_ray_tomography(a_hat_0, b):  
    "Pseudocode for carrying out X ray tomography"  
  
    # a_hat_0 is the initial image guess, of shape (NX, NY)  
    # b are the observed measurements, of shape (MX, MY)  
  
    a_hat = a_hat_0  
    lam = 1  
    for i in range(0, n_steps):  
        a_hat = a_hat.requires_grad_(True)  
        b_hat = numerical_integrate(a_hat)  
        R = total_variation(b_hat)  
        loss = torch.mean((b-b_hat)**2) + lam*R  
        da = torch.autograd.grad(loss, a_hat)  
        a_hat -= gamma*da  
  
    return a_hat
```

1. Start with initial guess \hat{a}

2. Loop:

1. Compute gradient, $\frac{\partial L(\hat{a})}{\partial \hat{a}}$

2. Take gradient descent step,

$$\hat{a} \leftarrow \hat{a} - \gamma \text{NN} \left(\frac{\partial L(\hat{a})}{\partial \hat{a}}, \hat{a}, \hat{b}, R(\hat{a}); \theta \right)$$

```
def Hybrid_X_ray_tomography(a_hat_0, b, theta):  
    "Pseudocode for carrying out X ray tomography, with NN correction"  
  
    # a_hat_0 is the initial image guess, of shape (NX, NY)  
    # b are the observed measurements, of shape (MX, MY)  
  
    a_hat = a_hat_0  
    lam = 1  
    for i in range(0, n_steps):  
        a_hat = a_hat.requires_grad_(True)  
        b_hat = numerical_integrate(a_hat)  
        R = total_variation(b_hat)  
        loss = torch.mean((b-b_hat)**2) + lam*R  
        da = torch.autograd.grad(loss, a_hat)  
        da = NN(da, a_hat, b_hat, R, theta)  
        a_hat -= gamma*da  
  
    return a_hat
```




Idea: **learn** a “better” direction to step in the parameter space


- How do we train this hybrid approach (learn θ)?

Hybrid X-ray tomography

Input to function:




Output:




```
def Hybrid_X_ray_tomography(a_hat_0, b, theta):  
    "Pseudocode for carrying out X ray tomography, with NN correction"  
  
    # a_hat_0 is the initial image guess, of shape (NX, NY)  
    # b are the observed measurements, of shape (MX, MY)  
  
    a_hat = a_hat_0  
    lam = 1  
    for i in range(0, n_steps):  
        a_hat = a_hat.requires_grad_(True)  
        b_hat = numerical_integrate(a_hat)  
        R = total_variation(b_hat)  
        loss = torch.mean((b-b_hat)**2) + lam*R  
        da = torch.autograd.grad(loss, a_hat)  
        da = NN(da, a_hat, b_hat, R, theta)  
        a_hat -= gamma*da  
  
    return a_hat
```

Hybrid X-ray tomography

Input to function:



Output:




```
def Hybrid_X_ray_tomography(a_hat_0, b, theta):  
    "Pseudocode for carrying out X ray tomography, with NN correction"  
  
    # a_hat_0 is the initial image guess, of shape (NX, NY)  
    # b are the observed measurements, of shape (MX, MY)  
  
    a_hat = a_hat_0  
    lam = 1  
    for i in range(0, n_steps):  
        a_hat = a_hat.requires_grad_(True)  
        b_hat = numerical_integrate(a_hat)  
        R = total_variation(b_hat)  
        loss = torch.mean((b-b_hat)**2) + lam*R  
        da = torch.autograd.grad(loss, a_hat)  
        da = NN(da, a_hat, b_hat, R, theta)  
        a_hat -= gamma*da  
  
    return a_hat
```

We train this hybrid approach using lots of examples of inputs/outputs (\hat{a}_0, b, a) and the loss function


$$L(\theta) = \sum_i^N \|H(\hat{a}_{0i}, b_i; \theta) - a_i\|^2$$

Hybrid X-ray tomography

Input to function:



Output:




We train this hybrid approach using lots of examples of inputs/outputs (\hat{a}_0, b, a) and the loss function

$$L(\theta) = \sum_i^N \|H(\hat{a}_{0i}, b_i; \theta) - a_i\|^2$$


```
def Hybrid_X_ray_tomography(a_hat_0, b, theta):  
    "Pseudocode for carrying out X ray tomography, with NN correction"  
  
    # a_hat_0 is the initial image guess, of shape (NX, NY)  
    # b are the observed measurements, of shape (MX, MY)  
  
    a_hat = a_hat_0  
    lam = 1  
    for i in range(0, n_steps):  
        a_hat = a_hat.requires_grad_(True)  
        b_hat = numerical_integrate(a_hat)  
        R = total_variation(b_hat)  
        loss = torch.mean((b-b_hat)**2) + lam*R  
        da = torch.autograd.grad(loss, a_hat)  
        da = NN(da, a_hat, b_hat, R, theta)  
        a_hat -= gamma*da  
  
    return a_hat  
  
# learn NN parameters  
theta.requires_grad_(True)  
for i in range(0, n_steps2):  
    a, b = # train NN using many example inverse problems  
    a_hat = Hybrid_X_ray_tomography(a_hat_0, b, theta)  
    loss = loss_fn(a, a_hat)  
    dtheta = torch.autograd.grad(loss, theta)  
    theta -= gamma*dtheta
```

Hybrid X-ray tomography

Input to function:



Output:



We train this hybrid approach using lots of examples of inputs/outputs (\hat{a}_0, b, a) and the loss function

$$L(\theta) = \sum_i^N \|H(\hat{a}_{0i}, b_i; \theta) - a_i\|^2$$


```
def Hybrid_X_ray_tomography(a_hat_0, b, theta):  
    "Pseudocode for carrying out X ray tomography, with NN correction"  
  
    # a_hat_0 is the initial image guess, of shape (NX, NY)  
    # b are the observed measurements, of shape (MX, MY)  
  
    a_hat = a_hat_0  
    lam = 1  
    for i in range(0, n_steps):  
        a_hat = a_hat.requires_grad_(True)  
        b_hat = numerical_integrate(a_hat)  
        R = total_variation(b_hat)  
        loss = torch.mean((b-b_hat)**2) + lam*R  
        da = torch.autograd.grad(loss, a_hat)  
        da = NN(da, a_hat, b_hat, R, theta)  
        a_hat -= gamma*da  
  
    return a_hat  
  
# learn NN parameters  
theta.requires_grad_(True)  
for i in range(0, n_steps2):  
    a, b = # train NN using many example inverse problems  
    a_hat = Hybrid_X_ray_tomography(a_hat_0, b, theta)  
    loss = loss_fn(a, a_hat)  
    dtheta = torch.autograd.grad(loss, theta)  
    theta -= gamma*dtheta
```

“Gradient descent on gradient descent”


“Learned gradient descent”

“Learning to learn”


Hybrid X-ray tomography



Ground truth



Traditional inversion



Learned gradient descent

Adler et al, Solving ill-posed inverse problems using iterative deep neural networks, Inverse Problems (2017)

Adding even more flexibility

- We can use **more** than one learnable component if we want!
- Where else would it be useful to add another?

```
def Hybrid_X_ray_tomography(a_hat_0, b, theta):  
    "Pseudocode for carrying out X ray tomography, with NN correction"  
  
    # a_hat_0 is the initial image guess, of shape (NX, NY)  
    # b are the observed measurements, of shape (MX, MY)  
  
    a_hat = a_hat_0  
    lam = 1  
    for i in range(0, n_steps):  
        a_hat.requires_grad_(True)  
        b_hat = numerical_integrate(a_hat)  
        R = total_variation(b_hat)  
        loss = torch.mean((b-b_hat)**2) + lam*R  
        da = torch.autograd.grad(loss, a_hat)  
        da = NN(da, a_hat, b_hat, R, theta)  
        a_hat -= gamma*da  
  
    return a_hat  
  
# learn NN parameters  
theta.requires_grad_(True)  
for i in range(0, n_steps2):  
    a, b = # train NN using many example inverse problems  
    a_hat = Hybrid_X_ray_tomography(a_hat_0, b, theta)  
    loss = loss_fn(a, a_hat)  
    dtheta = torch.autograd.grad(loss, theta)  
    theta -= gamma*dtheta
```

Idea: **learn** a “better” direction to step in the parameter space

Adding even more flexibility

```
def Hybrid2_X_ray_tomography(a_hat_0, b, theta):
    "Pseudocode for carrying out X ray tomography, with NN correction"

    # a_hat_0 is the initial image guess, of shape (NX, NY)
    # b are the observed measurements, of shape (MX, MY)

    a_hat = a_hat_0

    for i in range(0, n_steps):
        a_hat = a_hat.requires_grad_(True)
        b_hat = numerical_integrate(a_hat)
        R = total_variation(b_hat)
        loss = torch.mean((b-b_hat)**2) + theta[0]*R
        da = torch.autograd.grad(loss, a_hat)
        da = NN(da, a_hat, b_hat, R, theta[1])
        a_hat -= gamma*da

    return a_hat

# learn NN parameters
theta.requires_grad_(True)
for i in range(0, n_steps2):
    a, b = # train NN using many example inverse problems
    a_hat = Hybrid2_X_ray_tomography(a_hat_0, b, theta)
    loss = loss_fn(a, a_hat)
    dtheta = torch.autograd.grad(loss, theta)
    theta -= gamma*dtheta
```

Idea 2: learn regularisation **hyperparameter** too

```
def Hybrid_X_ray_tomography(a_hat_0, b, theta):
    "Pseudocode for carrying out X ray tomography, with NN correction"

    # a_hat_0 is the initial image guess, of shape (NX, NY)
    # b are the observed measurements, of shape (MX, MY)

    a_hat = a_hat_0
    lam = 1

    for i in range(0, n_steps):
        a_hat = a_hat.requires_grad_(True)
        b_hat = numerical_integrate(a_hat)
        R = total_variation(b_hat)
        loss = torch.mean((b-b_hat)**2) + lam*R
        da = torch.autograd.grad(loss, a_hat)
        da = NN(da, a_hat, b_hat, R, theta)
        a_hat -= gamma*da

    return a_hat

# learn NN parameters
theta.requires_grad_(True)
for i in range(0, n_steps2):
    a, b = # train NN using many example inverse problems
    a_hat = Hybrid_X_ray_tomography(a_hat_0, b, theta)
    loss = loss_fn(a, a_hat)
    dtheta = torch.autograd.grad(loss, theta)
    theta -= gamma*dtheta
```

Idea: **learn** a “better” direction to step in the parameter space

Adding even more flexibility

```
def Hybrid2_X_ray_tomography(a_hat_0, b, theta):
    "Pseudocode for carrying out X ray tomography, with NN correction"

    # a_hat_0 is the initial image guess, of shape (NX, NY)
    # b are the observed measurements, of shape (MX, MY)

    a_hat = a_hat_0

    for i in range(0, n_steps):
        a_hat = a_hat.requires_grad_(True)
        b_hat = numerical_integrate(a_hat)
        R = total_variation(b_hat)
        loss = torch.mean((b-b_hat)**2) + theta[0]*R
        da = torch.autograd.grad(loss, a_hat)
        da = NN(da, a_hat, b_hat, R, theta[1])
        a_hat -= gamma*da

    return a_hat

# learn NN parameters
theta.requires_grad_(True)
for i in range(0, n_steps2):
    a, b = # train NN using many example inverse problems
    a_hat = Hybrid2_X_ray_tomography(a_hat_0, b, theta)
    loss = loss_fn(a, a_hat)
    dtheta = torch.autograd.grad(loss, theta)
    theta -= gamma*dtheta
```

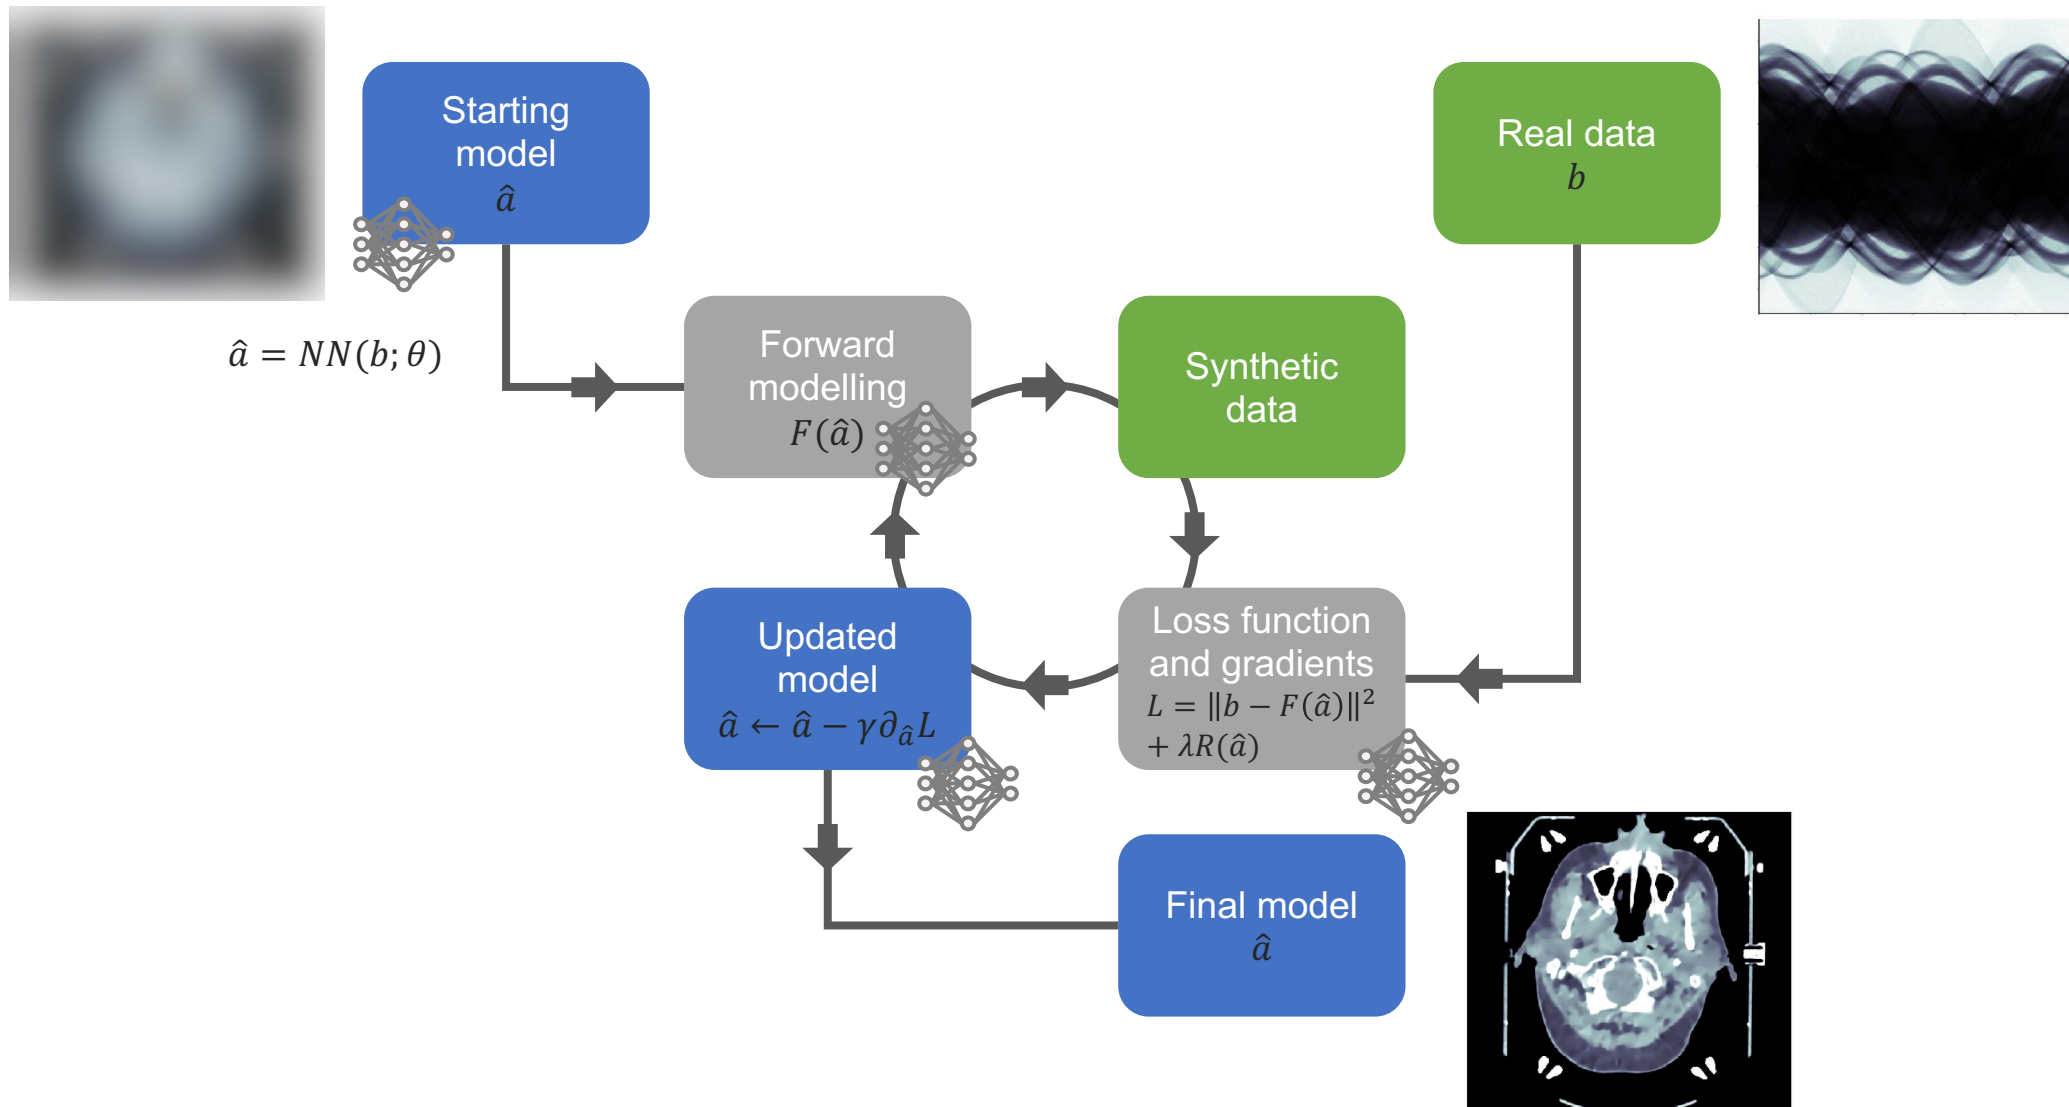


Key idea:

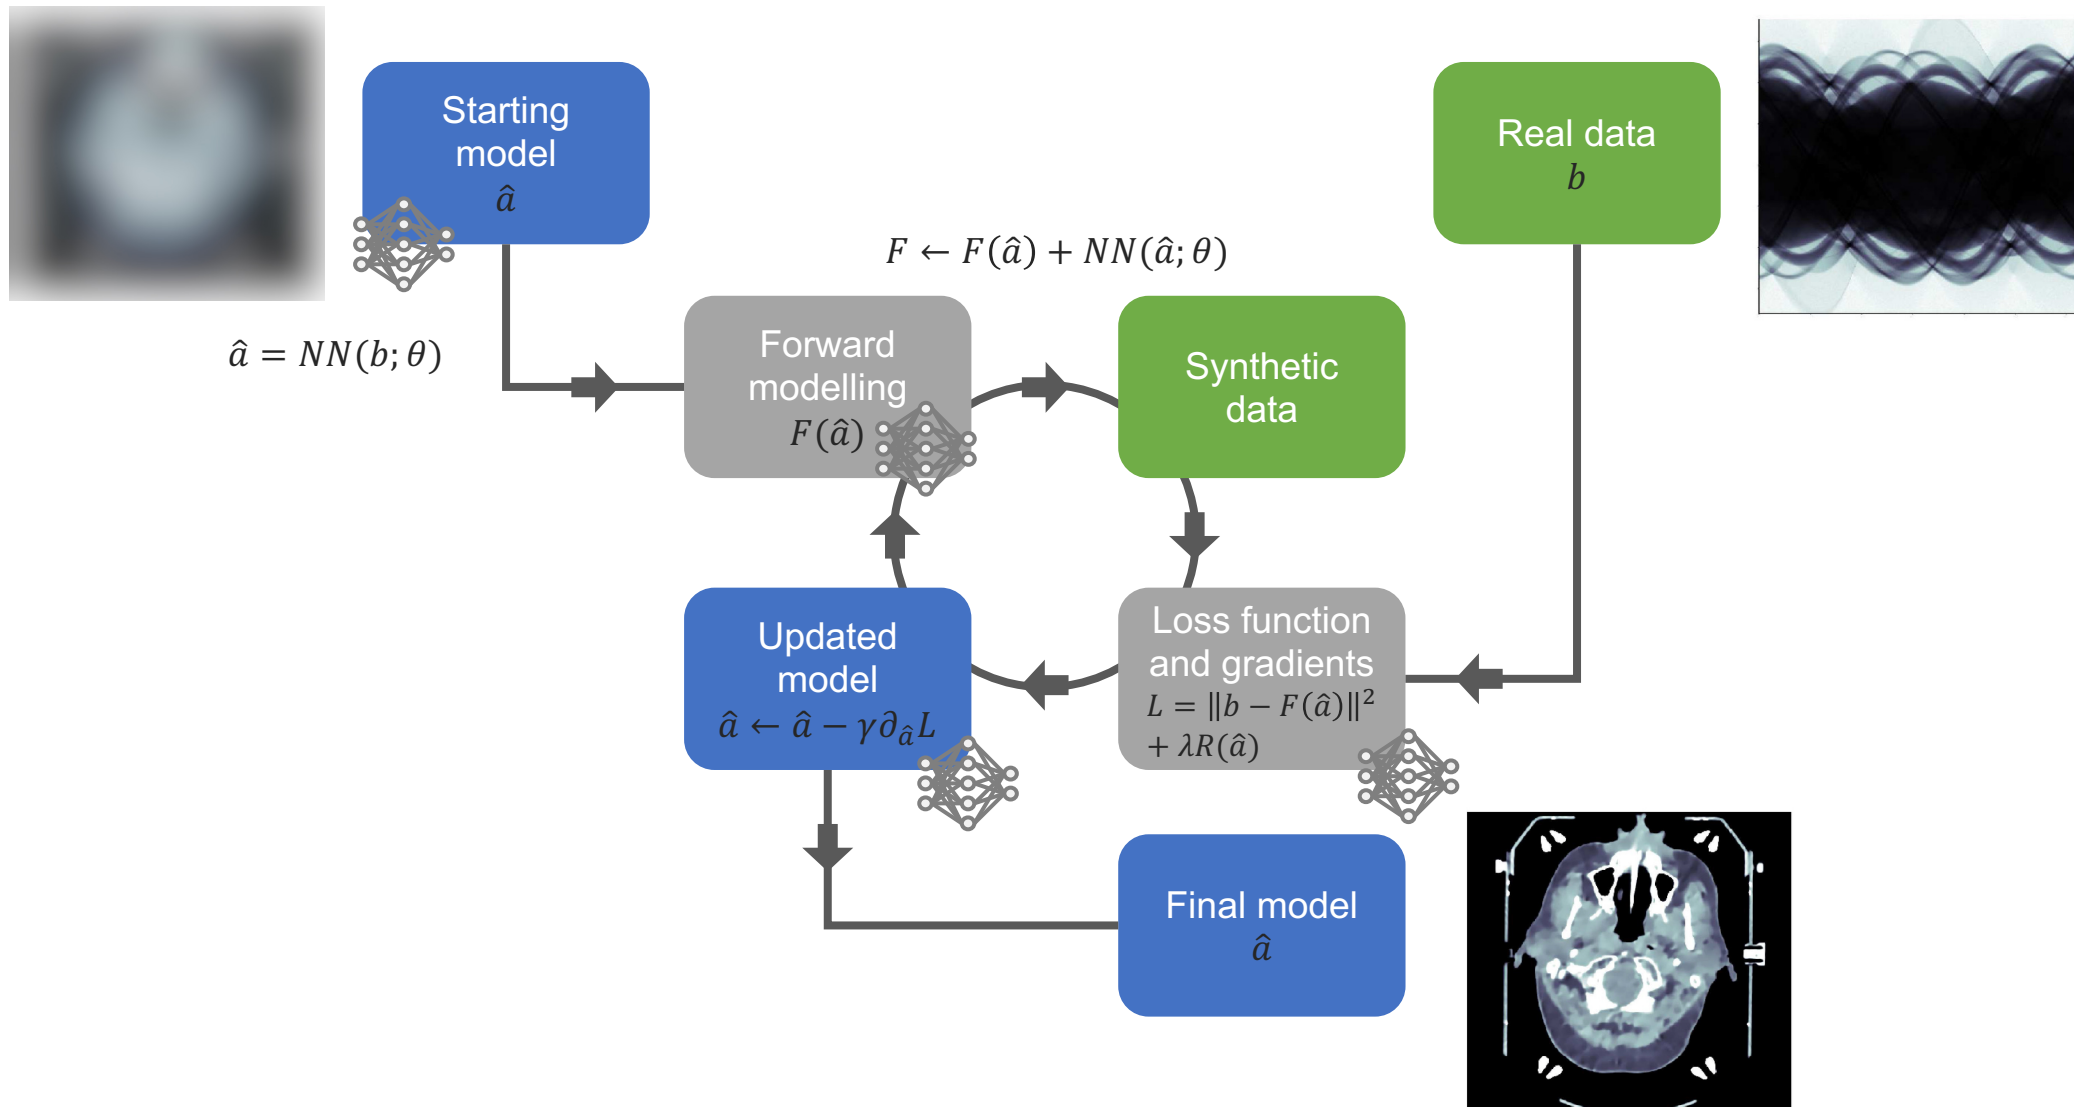
Traditional algorithms can be made as **learnable** (flexible) or as **unlearnable** (rigid) as you like

This allows you to balance the pros/cons of using NNs!

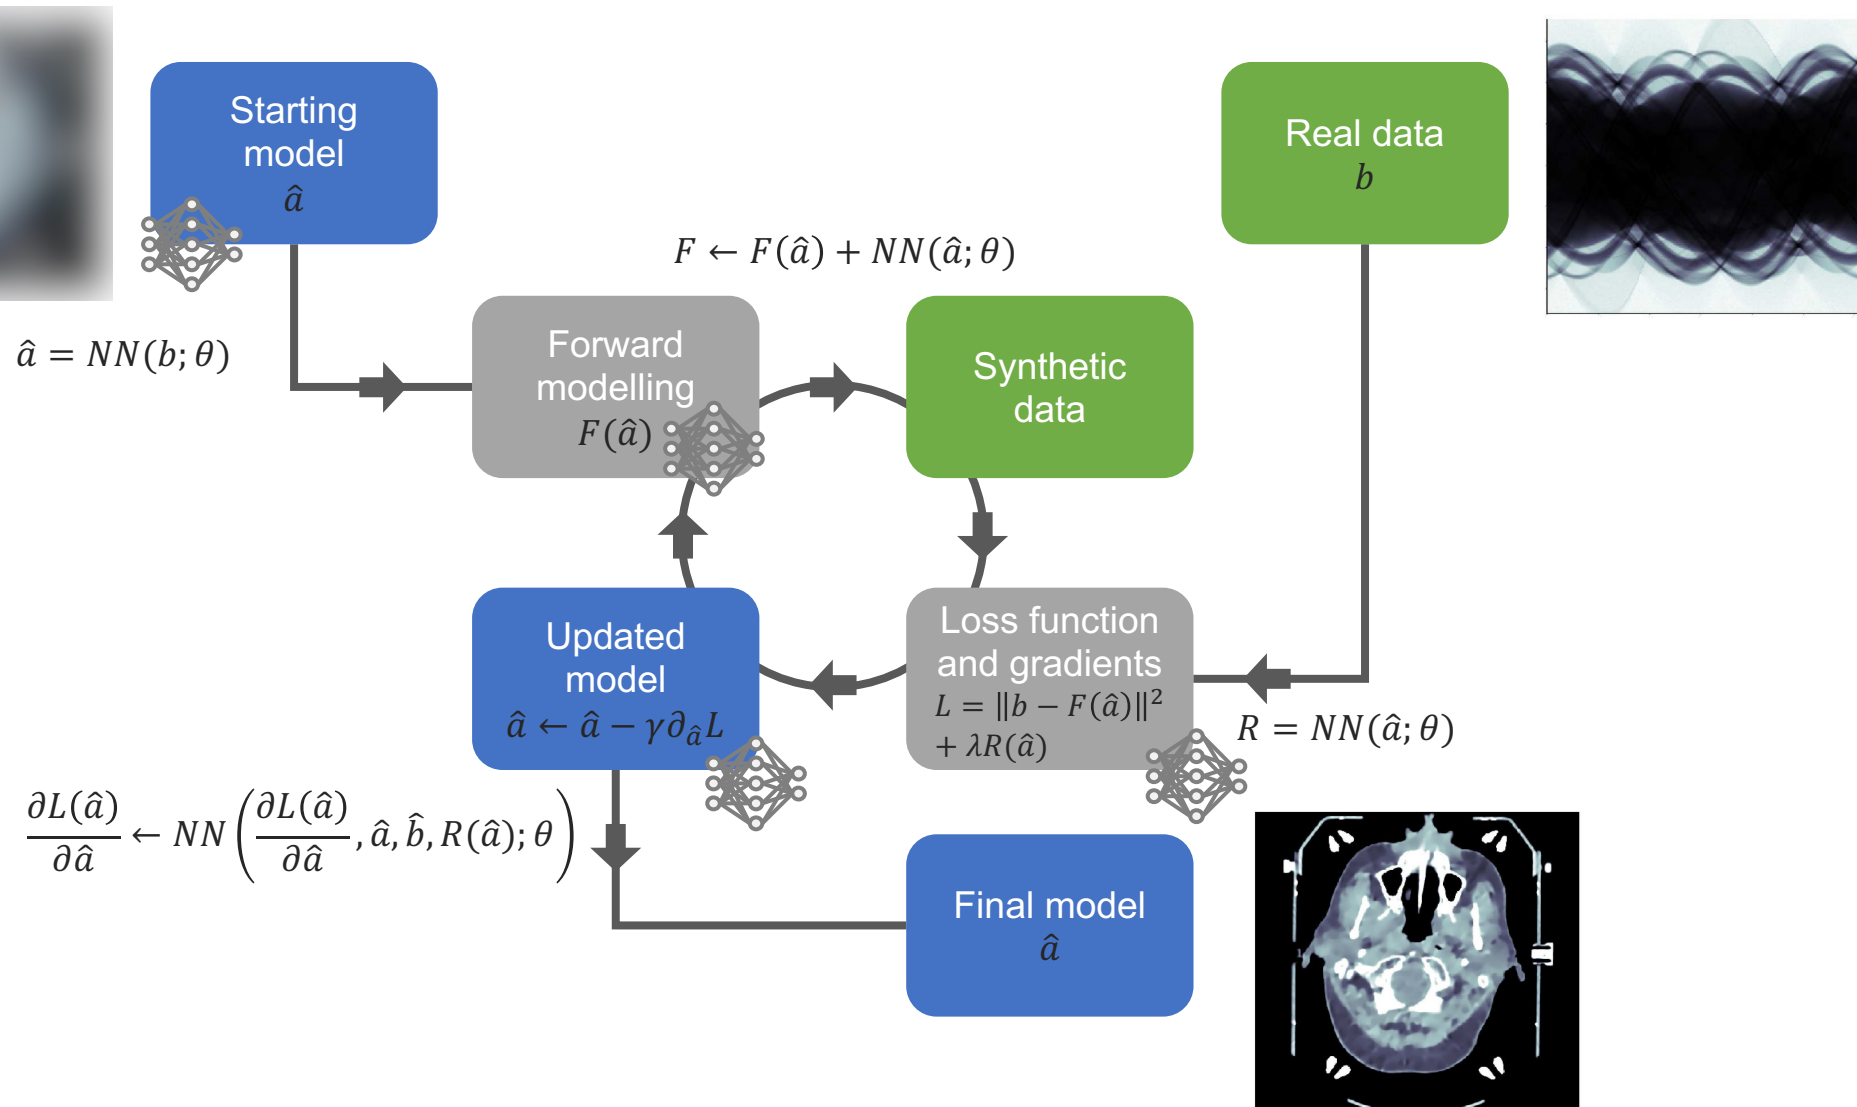
We can add learnable components everywhere!



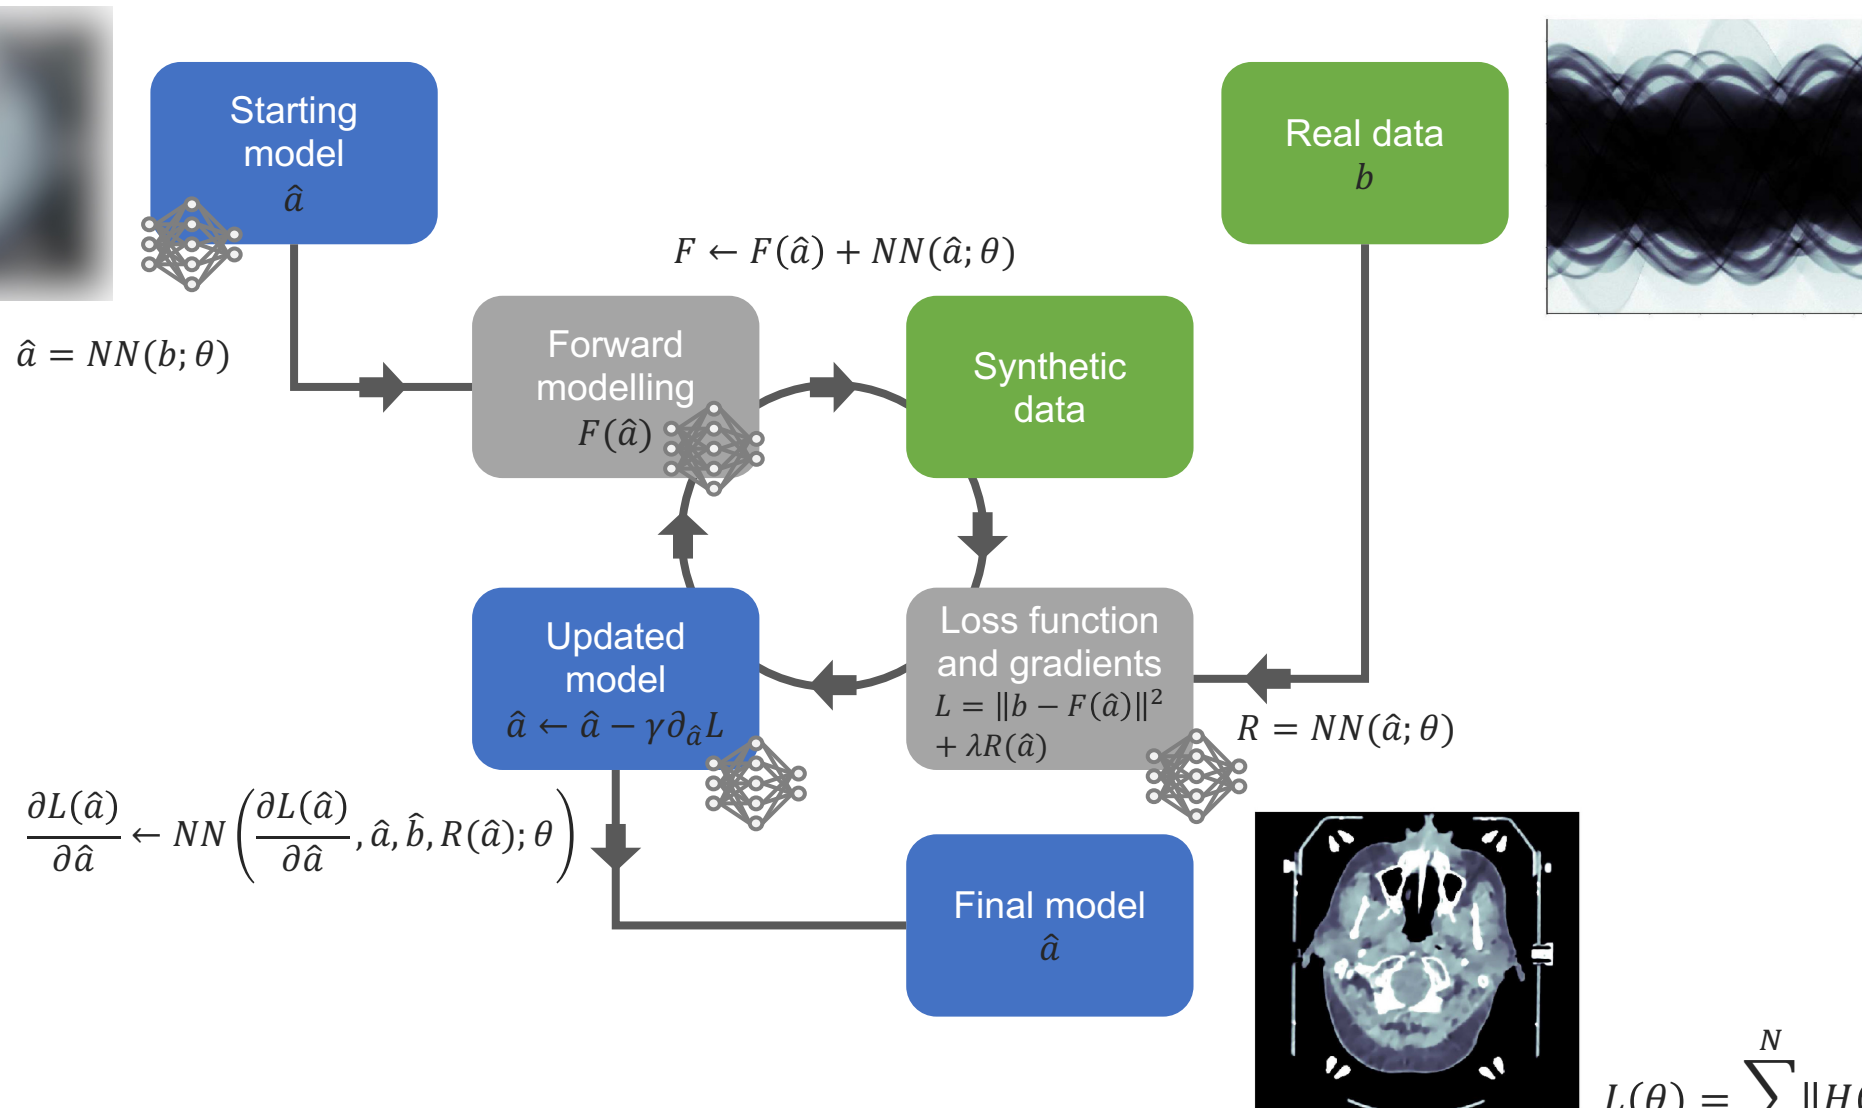
We can add learnable components everywhere!



We can add learnable components everywhere!



We can add learnable components everywhere!



$$L(\theta) = \sum_i^N \|H(\hat{a}_{0i}, b_i; \theta) - a_i\|^2$$

5 min break

Lecture overview

- Differentiable physics recap
- Coding a simple hybrid approach in PyTorch
- Hybrid approaches for inverse problems
- Neural differential equations (NDEs)
- Course summary

Ordinary differential equations

Consider solving an **ordinary differential equation** (ODE):

$$\begin{aligned}\frac{d\mathbf{x}(t)}{dt} &= \mathbf{f}(\mathbf{x}, t) \\ \mathbf{x}(t = 0) &= \mathbf{x}_0\end{aligned}$$

For example, the Lotka-Volterra system:


$$\begin{aligned}\frac{dx}{dt} &= \alpha x - \beta xy \\ \frac{dy}{dt} &= \gamma xy - \delta y\end{aligned}$$

x = population density of prey

y = population density of predator

α, β = max prey birth rate, effect of predators on prey growth rate

δ, γ = max predator death rate, effect of prey on predator growth rate



Source: wikipedia

$$\alpha, \beta = 1.1, 0.4$$

$$\delta, \gamma = 0.4, 0.1$$

$$x_0 = y_0 = 10$$


Ordinary differential equations

Consider solving an **ordinary differential equation** (ODE):

$$\begin{aligned}\frac{d\mathbf{x}(t)}{dt} &= \mathbf{f}(\mathbf{x}, t) \\ \mathbf{x}(t=0) &= \mathbf{x}_0\end{aligned}$$

For example, the Lotka-Volterra system:

$$\begin{aligned}\frac{dx}{dt} &= \alpha x - \beta xy \\ \frac{dy}{dt} &= \gamma xy - \delta y\end{aligned}$$



Source: wikipedia

- What if we are unsure of $\mathbf{f}(\mathbf{x}, t)$?


Neural differential equations

Consider solving an **ordinary differential equation** (ODE):

$$\begin{aligned}\frac{dx(t)}{dt} &= f(x, t; \theta) \\ x(t=0) &= x_0\end{aligned}$$

For example, the Lotka-Volterra system:

$$\begin{aligned}\frac{dx}{dt} &= \alpha x + NN(x, y; \theta_1) \\ \frac{dy}{dt} &= NN(x, y; \theta_2) - \theta_3 y\end{aligned}$$



Source: wikipedia

- What if we are unsure of $f(x, t)$?
- Use neural networks to represent uncertain parts


Neural differential equations

Consider solving an **ordinary differential equation** (ODE):

$$\begin{aligned}\frac{dx(t)}{dt} &= f(x, t; \theta) \\ x(t = 0) &= x_0\end{aligned}$$

For example, the Lotka-Volterra system:

$$\begin{aligned}\frac{dx}{dt} &= \alpha x + NN(x, y; \theta_1) \\ \frac{dy}{dt} &= NN(x, y; \theta_2) - \theta_3 y\end{aligned}$$



Source: wikipedia

- What if we are unsure of $f(x, t)$?
- Use neural networks to represent uncertain parts
- This is known as a **neural differential equation**


Neural differential equations

Consider solving an **ordinary differential equation** (ODE):

$$\begin{aligned}\frac{dx(t)}{dt} &= f(x, t; \theta) \\ x(t=0) &= x_0\end{aligned}$$

For example, the Lotka-Volterra system:

$$\begin{aligned}\frac{dx}{dt} &= \alpha x + NN(x, y; \theta_1) \\ \frac{dy}{dt} &= NN(x, y; \theta_2) - \theta_3 y\end{aligned}$$



Source: wikipedia

- What if we are unsure of $f(x, t)$?
- Use neural networks to represent uncertain parts
- This is known as a **neural differential equation**
- We can use a hybrid approach to **learn** the dynamics

Neural differential equations

Consider solving an **ordinary differential equation** (ODE):

$$\begin{aligned}\frac{dx(t)}{dt} &= f(x, t; \theta) \\ x(t = 0) &= x_0\end{aligned}$$

For example, the Lotka-Volterra system:

$$\begin{aligned}\frac{dx}{dt} &= \alpha x + NN(x, y; \theta_1) \\ \frac{dy}{dt} &= NN(x, y; \theta_2) - \theta_3 y\end{aligned}$$

- What if we are unsure of $f(x, t)$?
- Use neural networks to represent uncertain parts
- This is known as a **neural differential equation**
- We can use a hybrid approach to **learn** the dynamics

A simple way of solving an ODE is to **discretise** in time and use the **Euler method**:

Given $x_0, t_0, \delta t$:

$$\begin{aligned}x_{i+1} &= x_i + \delta t f(x_i, t_i) \\ t_{i+1} &= t_i + \delta t\end{aligned}$$

For the Lotka-Volterra system:

$$\begin{aligned}x_{i+1} &= x_i + \delta t(\alpha x_i - \beta x_i y_i) \\ y_{i+1} &= y_i + \delta t(\gamma x_i y_i - \delta y_i) \\ t_{i+1} &= t_i + \delta t\end{aligned}$$

For the learnable Lotka-Volterra system:

$$\begin{aligned}x_{i+1} &= x_i + \delta t(\alpha x_i + NN(x_i, y_i; \theta_1)) \\ y_{i+1} &= y_i + \delta t(NN(x_i, y_i; \theta_2) - \theta_3 y_i) \\ t_{i+1} &= t_i + \delta t\end{aligned}$$

Hybrid Lotka-Volterra solver

A simple way of solving an ODE is to **discretise** in time and use the **Euler method**:

Given $x_0, t_0, \delta t$:

$$\begin{aligned}x_{i+1} &= x_i + \delta t f(x_i, t_i) \\ t_{i+1} &= t_i + \delta t\end{aligned}$$

For the Lotka-Volterra system:

$$\begin{aligned}x_{i+1} &= x_i + \delta t(\alpha x_i - \beta x_i y_i) \\ y_{i+1} &= y_i + \delta t(\gamma x_i y_i - \delta y_i) \\ t_{i+1} &= t_i + \delta t\end{aligned}$$

For the learnable Lotka-Volterra system:

$$\begin{aligned}x_{i+1} &= x_i + \delta t(\alpha x_i + NN(x_i, y_i; \theta_1)) \\ y_{i+1} &= y_i + \delta t(NN(x_i, y_i; \theta_2) - \theta_3 y_i) \\ t_{i+1} &= t_i + \delta t\end{aligned}$$


```
def Hybrid_LV_Euler_solver(x0, y0, dt, theta):  
    """Pseudocode for solving Lotka-Volterra system,  
    with learnable dynamics"""  
  
    x, y = x0, y0  
    for t in range(0, T):  
        x = x + dt*(alpha*x + NN(x, y, theta[0]))  
        y = y + dt*(NN(x, y, theta[1]) - theta[2]*y)  
    return x, y
```

Hybrid Lotka-Volterra solver

Train the hybrid solver using loss function:

$$L(\theta) = \sum_i^T \|x_i(x_0, t_0, \delta t, \theta) - x_{\text{observed } i}\|^2$$


```
def Hybrid_LV_Euler_solver(x0, y0, dt, theta):  
    """Pseudocode for solving Lotka-Volterra system,  
    with learnable dynamics"""  
  
    x, y = x0, y0  
    for t in range(0, T):  
        x = x + dt*(alpha*x + NN(x, y, theta[0]))  
        y = y + dt*(NN(x, y, theta[1]) - theta[2]*y)  
    return x, y
```



Rackauckas et al, Universal differential equations for scientific machine learning, ArXiv (2021)


Hybrid Lotka-Volterra solver

Rackauckas et al, Universal differential equations for scientific machine learning, ArXiv (2021)



Hybrid Lotka-Volterra solver

Rackauckas et al, Universal differential equations for scientific machine learning, ArXiv (2021)




$$\begin{aligned}\frac{dx}{dt} &= \alpha x + NN(x, y; \theta_1) \\ \frac{dy}{dt} &= NN(x, y; \theta_2) - \theta_3 y\end{aligned}$$

- Note, after training, we can do **symbolic regression** on $NN(x, y; \theta_1)$ and $NN(x, y; \theta_2)$ to “discover” their functional form, i.e. that $NN(x, y; \theta_1) \approx -\beta xy$

Hybrid Lotka-Volterra solver

Rackauckas et al, Universal differential equations for scientific machine learning, ArXiv (2021)




$$\begin{aligned}\frac{dx}{dt} &= \alpha x + NN(x, y; \theta_1) \\ \frac{dy}{dt} &= NN(x, y; \theta_2) - \theta_3 y\end{aligned}$$

- This model generalizes well!

Hybrid Lotka-Volterra solver

Rackauckas et al, Universal differential equations for scientific machine learning, ArXiv (2021)




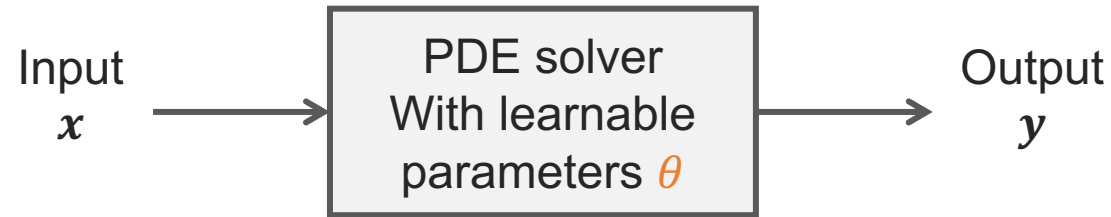
$$\frac{dx}{dt} = \alpha x + NN(x, y; \theta_1)$$
$$\frac{dy}{dt} = NN(x, y; \theta_2) - \theta_3 y$$

- This model generalizes well!



Idea: what if we use differential equations to model **any** dataset (not just physical systems)?

Neural differential equations


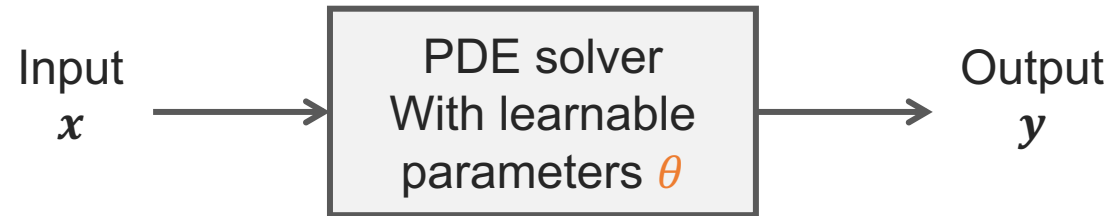


$$P(x = 7)$$



Idea: what if we use differential equations to model **any** dataset (not just physical systems)?

Neural differential equations



$$P(x = 7)$$



Idea: what if we use differential equations to model **any** dataset (not just physical systems)?

- We can think of the PDE solver as a “custom” NN architecture

Neural ordinary differential equations

Consider solving an **ordinary differential equation** (ODE):

$$\begin{aligned}\frac{dx(t)}{dt} &= f(x; \theta) \\ x(t=0) &= x_0\end{aligned}$$

Solver using Euler method:

Given $x_0, \delta t$:

$$x_{i+1} = x_i + \delta t f(x_i; \theta)$$

Neural ordinary differential equations


Consider solving an **ordinary differential equation** (ODE):

$$\frac{dx(t)}{dt} = f(x; \theta)$$
$$x(t = 0) = x_0$$

Solver using Euler method:

Given $x_0, \delta t$:

$$x_{i+1} = x_i + \delta t f(x_i; \theta)$$



Neural ordinary differential equations


Consider solving an **ordinary differential equation** (ODE):

$$\begin{aligned}\frac{dx(t)}{dt} &= f(x; \theta) \\ x(t=0) &= x_0\end{aligned}$$

Solver using Euler method:

Given $x_0, \delta t$:

$$x_{i+1} = x_i + \delta t f(x_i; \theta)$$




This is **identical** to the **residual layer** used in standard residual networks (ResNets)!

ResNets are Euler solvers

I.e., ResNets are Euler solvers!

⇒ In the **limit** of an infinite numbers of layers (i.e. as $\delta t \rightarrow 0$), a ResNet computes the solution to

$$\frac{dx(t)}{dt} = f(x, \theta)$$



This is **identical** to the **residual layer** used in standard residual networks (ResNets)!

Higher-order solvers

Consider solving an **ordinary differential equation** (ODE):

$$\begin{aligned}\frac{dx(t)}{dt} &= f(x; \theta) \\ x(t=0) &= x_0\end{aligned}$$

We are not limited to Euler solvers! Many **other** solvers could be used, for example higher-order Runge-Kutta methods, e.g. RK4:

$$x_{i+1} = x_i + \frac{\delta t}{6} (k_1 + 2k_2 + 2k_3 + k_4)$$

$$k_1 = f(x_i; \theta)$$

$$k_2 = f\left(x_i + \frac{\delta t}{2} k_1; \theta\right)$$

$$k_3 = f\left(x_i + \frac{\delta t}{2} k_2; \theta\right)$$

$$k_4 = f(x_i + \delta t k_3; \theta)$$

Higher-order solvers

Consider solving an **ordinary differential equation** (ODE):

$$\frac{dx(t)}{dt} = f(x; \theta)$$
$$x(t = 0) = x_0$$

We are not limited to Euler solvers! Many **other** solvers could be used, for example higher-order Runge-Kutta methods, e.g. RK4:

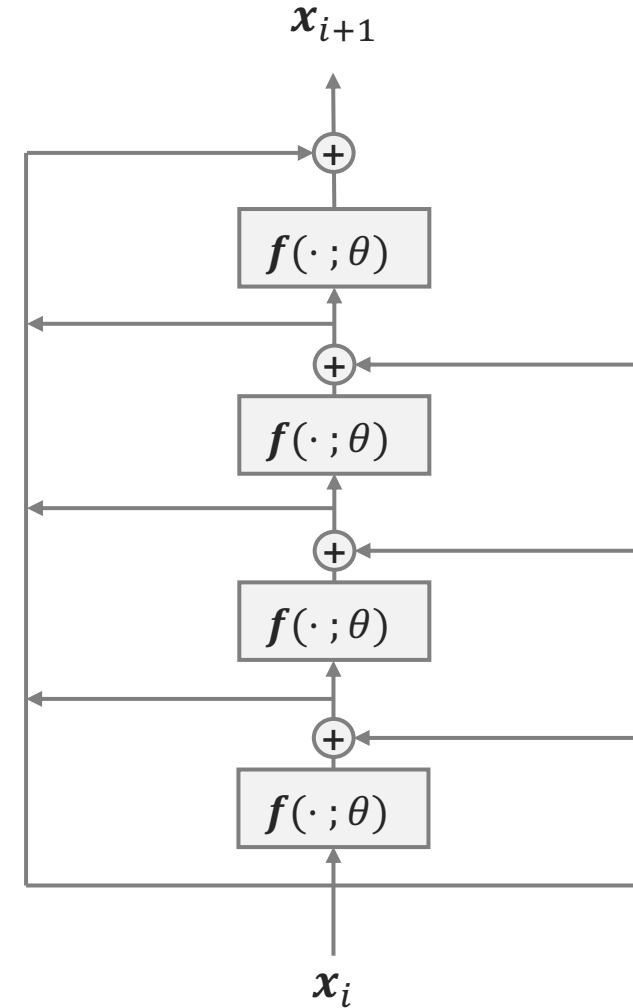
$$x_{i+1} = x_i + \frac{\delta t}{6} (k_1 + 2k_2 + 2k_3 + k_4)$$

$$k_1 = f(x_i; \theta)$$

$$k_2 = f\left(x_i + \frac{\delta t}{2} k_1; \theta\right)$$

$$k_3 = f\left(x_i + \frac{\delta t}{2} k_2; \theta\right)$$

$$k_4 = f(x_i + \delta t k_3; \theta)$$




Higher-order solvers

Consider solving an **ordinary differential equation** (ODE):

$$\begin{aligned}\frac{dx(t)}{dt} &= f(x; \theta) \\ x(t=0) &= x_0\end{aligned}$$

We are not limited to Euler solvers! Many **other** solvers could be used, for example higher-order Runge-Kutta methods, e.g. RK4:

“Custom” residual block



Higher-order solvers


Consider solving an **ordinary differential equation** (ODE):

$$\frac{dx(t)}{dt} = f(x; \theta)$$
$$x(t = 0) = x_0$$

We are not limited to Euler solvers! Many **other** solvers could be used, for example higher-order Runge-Kutta methods, e.g. RK4:

=> Other solvers define other NN architectures

“Custom” residual block



Higher-order solvers

Consider solving an **ordinary differential equation** (ODE):

$$\frac{dx(t)}{dt} = f(x; \theta)$$
$$x(t = 0) = x_0$$

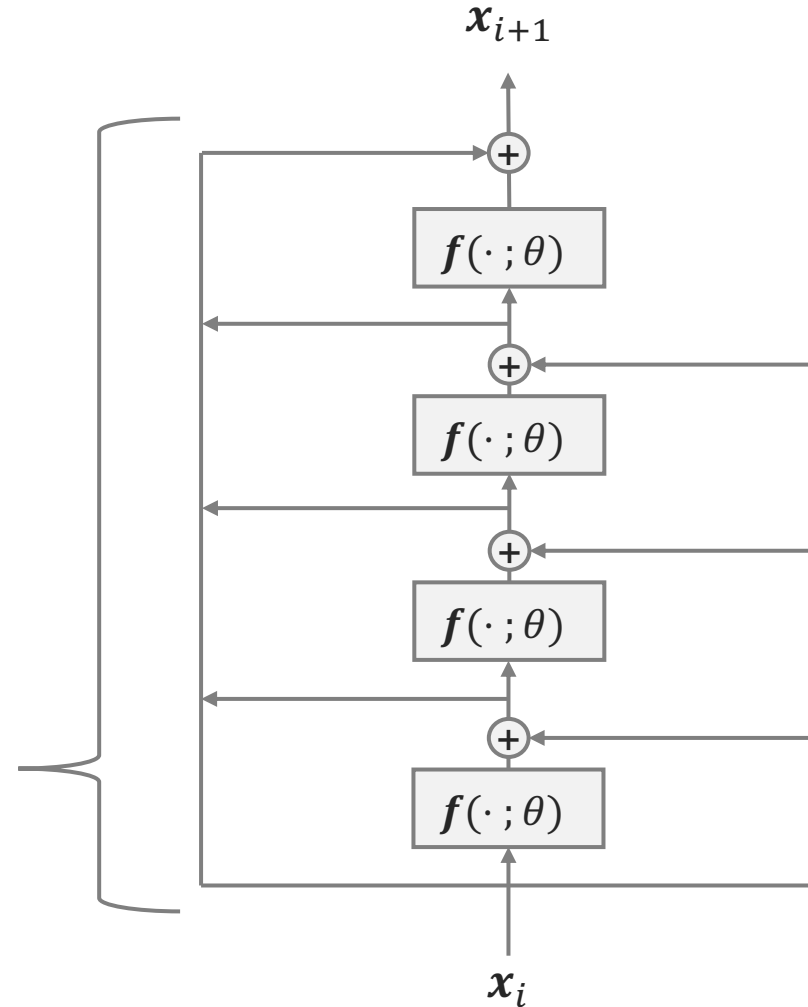
We are not limited to Euler solvers! Many **other** solvers could be used, for example higher-order Runge-Kutta methods, e.g. RK4:

	Test Error
1-Layer MLP [†]	1.60%
ResNet	0.41%
RK-Net	0.47%

Performance on MNIST (digit classification)

Chen et al, Neural ordinary differential equations, NeurIPS (2018)


“Custom” residual block



Neural differential equations

Consider a 1D convolutional layer:

$$\begin{aligned} y_{i+1} &= \theta \star y_i \\ &= (\theta_1 \quad \theta_2 \quad \theta_3) \star y_i \end{aligned}$$



Neural differential equations


Consider a 1D convolutional layer:

$$\begin{aligned} y_{i+1} &= \theta \star y_i \\ &= (\theta_1 \quad \theta_2 \quad \theta_3) \star y_i \end{aligned}$$

Let us **transform** θ to a new vector $\beta(\theta)$ which is (uniquely) given by

$$\begin{pmatrix} \frac{1}{4} & -\frac{1}{2h} & -\frac{1}{h^2} \\ \frac{1}{2} & 0 & \frac{2}{h^2} \\ \frac{1}{4} & \frac{1}{2h} & -\frac{1}{h^2} \end{pmatrix} \begin{pmatrix} \beta_1 \\ \beta_2 \\ \beta_3 \end{pmatrix} = \begin{pmatrix} \theta_1 \\ \theta_2 \\ \theta_3 \end{pmatrix}$$


For some $h > 0$.



Neural differential equations

Consider a 1D convolutional layer:

$$\begin{aligned} y_{i+1} &= \theta \star y_i \\ &= (\theta_1 \quad \theta_2 \quad \theta_3) \star y_i \end{aligned}$$



Let us **transform** θ to a new vector $\beta(\theta)$ which is (uniquely) given by

$$\begin{pmatrix} \frac{1}{4} & -\frac{1}{2h} & -\frac{1}{h^2} \\ \frac{1}{2} & 0 & \frac{2}{h^2} \\ \frac{1}{4} & \frac{1}{2h} & -\frac{1}{h^2} \end{pmatrix} \begin{pmatrix} \beta_1 \\ \beta_2 \\ \beta_3 \end{pmatrix} = \begin{pmatrix} \theta_1 \\ \theta_2 \\ \theta_3 \end{pmatrix}$$

For some $h > 0$.

Then we can **re-write** the convolutional layer as


$$y_{i+1} = \left(\frac{\beta_1(\theta)}{4} (1 \quad 2 \quad 1) + \frac{\beta_2(\theta)}{2h} (-1 \quad 0 \quad 1) + \frac{\beta_3(\theta)}{h^2} (-1 \quad 2 \quad -1) \right) \star y_i$$

Ruthotto and Haber, Deep Neural Networks Motivated by Partial
Differential Equations, Journal of Mathematical Imaging and Vision (2019)

Neural differential equations

Consider a 1D convolutional layer:

$$\begin{aligned} y_{i+1} &= \theta \star y_i \\ &= (\theta_1 \quad \theta_2 \quad \theta_3) \star y_i \end{aligned}$$



Let us **transform** θ to a new vector $\beta(\theta)$ which is (uniquely) given by

In the **limit** $h \rightarrow 0$,

$$\begin{pmatrix} \frac{1}{4} & -\frac{1}{2h} & -\frac{1}{h^2} \\ \frac{1}{2} & 0 & \frac{2}{h^2} \\ \frac{1}{4} & \frac{1}{2h} & -\frac{1}{h^2} \end{pmatrix} \begin{pmatrix} \beta_1 \\ \beta_2 \\ \beta_3 \end{pmatrix} = \begin{pmatrix} \theta_1 \\ \theta_2 \\ \theta_3 \end{pmatrix}$$

$$y_{i+1} = \beta_1(\theta)y_i + \beta_2(\theta)\frac{\partial y_i}{\partial x} + \beta_3(\theta)\frac{\partial^2 y_i}{\partial x^2}$$

For some $h > 0$.


Then we can **re-write** the convolutional layer as

$$y_{i+1} = \left(\frac{\beta_1(\theta)}{4} (1 \quad 2 \quad 1) + \frac{\beta_2(\theta)}{2h} (-1 \quad 0 \quad 1) + \frac{\beta_3(\theta)}{h^2} (-1 \quad 2 \quad -1) \right) \star y_i$$

Neural differential equations

Consider a 1D convolutional layer:

$$\begin{aligned} y_{i+1} &= \theta \star y_i \\ &= (\theta_1 \quad \theta_2 \quad \theta_3) \star y_i \end{aligned}$$



Let us **transform** θ to a new vector $\beta(\theta)$ which is (uniquely) given by

$$\begin{pmatrix} \frac{1}{4} & -\frac{1}{2h} & -\frac{1}{h^2} \\ \frac{1}{2} & 0 & \frac{2}{h^2} \\ \frac{1}{4} & \frac{1}{2h} & -\frac{1}{h^2} \end{pmatrix} \begin{pmatrix} \beta_1 \\ \beta_2 \\ \beta_3 \end{pmatrix} = \begin{pmatrix} \theta_1 \\ \theta_2 \\ \theta_3 \end{pmatrix}$$

For some $h > 0$.

Then we can **re-write** the convolutional layer as

$$y_{i+1} = \left(\frac{\beta_1(\theta)}{4} (1 \quad 2 \quad 1) + \frac{\beta_2(\theta)}{2h} (-1 \quad 0 \quad 1) + \frac{\beta_3(\theta)}{h^2} (-1 \quad 2 \quad -1) \right) \star y_i$$

In the **limit** $h \rightarrow 0$,

$$y_{i+1} = \beta_1(\theta) y_i + \beta_2(\theta) \frac{\partial y_i}{\partial x} + \beta_3(\theta) \frac{\partial^2 y_i}{\partial x^2}$$


Consider a residual CNN, then

$$y_{i+1} = y_i + \beta_1(\theta) y_i + \beta_2(\theta) \frac{\partial y_i}{\partial x} + \beta_3(\theta) \frac{\partial^2 y_i}{\partial x^2}$$

Neural differential equations

Consider a 1D convolutional layer:

$$\begin{aligned} y_{i+1} &= \theta \star y_i \\ &= (\theta_1 \quad \theta_2 \quad \theta_3) \star y_i \end{aligned}$$



Let us **transform** θ to a new vector $\beta(\theta)$ which is (uniquely) given by

$$\begin{pmatrix} \frac{1}{4} & -\frac{1}{2h} & -\frac{1}{h^2} \\ \frac{1}{2} & 0 & \frac{2}{h^2} \\ \frac{1}{4} & \frac{1}{2h} & -\frac{1}{h^2} \end{pmatrix} \begin{pmatrix} \beta_1 \\ \beta_2 \\ \beta_3 \end{pmatrix} = \begin{pmatrix} \theta_1 \\ \theta_2 \\ \theta_3 \end{pmatrix}$$

For some $h > 0$.

Then we can **re-write** the convolutional layer as

$$y_{i+1} = \left(\frac{\beta_1(\theta)}{4} (1 \quad 2 \quad 1) + \frac{\beta_2(\theta)}{2h} (-1 \quad 0 \quad 1) + \frac{\beta_3(\theta)}{h^2} (-1 \quad 2 \quad -1) \right) \star y_i$$

In the **limit** $h \rightarrow 0$,

$$y_{i+1} = \beta_1(\theta) y_i + \beta_2(\theta) \frac{\partial y_i}{\partial x} + \beta_3(\theta) \frac{\partial^2 y_i}{\partial x^2}$$

Consider a residual CNN, then

$$y_{i+1} = y_i + \beta_1(\theta) y_i + \beta_2(\theta) \frac{\partial y_i}{\partial x} + \beta_3(\theta) \frac{\partial^2 y_i}{\partial x^2}$$

In the **limit** of infinite layers, the residual CNN solves

$$\frac{\partial y}{\partial t} = \beta_1(\theta) y + \beta_2(\theta) \frac{\partial y}{\partial x} + \beta_3(\theta) \frac{\partial^2 y}{\partial x^2}$$

Neural differential equations

Neural network architectures \Leftrightarrow Differential equation solvers

Understanding of architectures / training algorithms \Leftrightarrow Understanding of PDEs / their solutions

Lecture overview

- Differentiable physics recap
- Coding a simple hybrid approach in PyTorch
- Hybrid approaches for inverse problems
- Neural differential equations (NDEs)
- Course summary

Scientific machine learning (SciML)

Major problem

Despite big breakthroughs in science + AI

Naively using deep learning for scientific tasks usually leads to:


- Lack of interpretability
- Poor generalisation
- Lots of training data required

Do neural networks really “**understand**” the scientific tasks they are being applied to?

Traditional scientific method:

- Revolves around theory and experiment
- a good theory should be explainable and make **novel** predictions

Solution



General trends in SciML

- Incorporating scientific understanding nearly always **improves** the performance of ML algorithms
- SciML approaches can be as flexible (learnable) or as inflexible (unlearnable) as necessary
- There are a plethora of SciML approaches; chose the one which **suits** your problem
- SciML approaches **still** suffer from the limitations of deep neural networks (generalisation, lack of interpretability, scalability, ...)
- SciML approaches can be applied to:
 - **many** different problems (simulation, inversion, data assimilation, control, equation discovery, ...)
 - **many** different fields
- SciML requires truly **interdisciplinary** research

Course learning objectives

- Aware of advanced **applications** of deep learning in scientific computing
- Familiar with the **design, implementation** and **theory** of these algorithms
- Understand the **pros/cons** of using deep learning
- Understand key scientific machine learning **concepts** and themes

Thank you!